

ECE 350

Real-time

Operating

Systems



Lecture I: Introduction

Prof. Seyed Majid Zahedi

<https://ece.uwaterloo.ca/~smzahedi>

Outline

- How do things work in ECE 350?
- What is a real-time system?
- What is an operating system?
- What makes operating systems so exciting?

Useful Links

- Course webpage

<https://ece.uwaterloo.ca/~smzahedi/crs/ece350>

- Course on Piazza

<https://piazza.com/uwaterloo.ca/winter2022/ece350>

- Anonymous feedback form

<https://forms.gle/c5kK1twHfRejuwMv9>

Class is Entirely Online!

- Lectures will be delivered on Teams
 - Links are provided on course webpage
 - Recordings will be available afterwards
- Office hours will be on Teams
 - Links are provided on course webpage
 - There will be no recordings
 - No office hours during week 1 and reading week
- Lab tutorials will be delivered on Teams
 - Links are provided on course webpage



Who Are We?



Instructor: Prof. Seyed Majid Zahedi
zahedi@uwaterloo.ca
<https://ece.uwaterloo.ca/~smzahedi>

Lab instructor: Aravind Vellora Vayalapra
avelloravayalapra@uwaterloo.ca



ECE 350 GTAs



Khilav Divyang Soni
k7soni@uwaterloo.ca

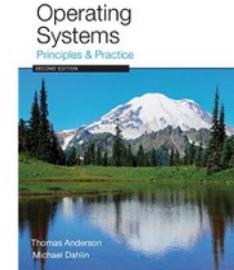
TBD
tbd@uwaterloo.ca



Readings

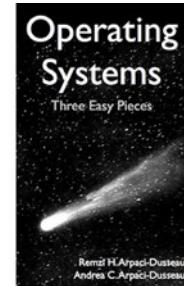
- Main textbook

[Operating Systems: Principles and Practice \(2nd Edition\)](#)

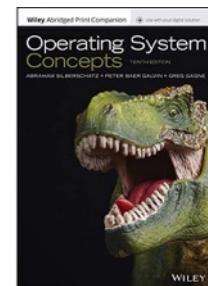


- Optional references

[Operating Systems: Three Easy Pieces \(Freely Available\)](#)



[Operating System Concepts \(10th Edition\)](#)



Prerequisite: ECE 252

- It is assumed that you know
 - File system API
 - File operations, directory structure
 - Processes and threads
 - PCB, TCB, process/thread life cycle, fork and exec, signals
 - Inter-process communication
 - File sharing, memory sharing, message passing, pipes
 - Networking
 - Sockets, ports, client and server programs
 - Concurrency
 - Mutual exclusion, mutex, semaphores, locks, condition variables, monitors
 - Deadlock
 - Avoidance, detection, recovery, Banker's algorithm

ECE 350 Is a Class About...

- Design of key systems abstractions that have emerged over time
 - Processes, threads, events, address spaces, file systems, sockets, transactions, key-value stores, etc.
- Tradeoffs surrounding these designs
- Their efficient implementation
 - Including hardware support that makes them possible and practical
- And how to use them effectively

Why Take ECE 350?

Why Learn About OS?

- Some of you will design and build parts of (real-time) operating systems
- Many of you will create systems that use OS concepts
 - Whether you build hardware or software
 - Concepts and design patterns appear at many levels
- All of you will write programs that use OS abstractions
 - The better you understand them, the better you use them

PAY ATTENTION!



**THIS IS THE IMPORTANT STUFF.
SERIOUSLY GUYS.**

Evaluation

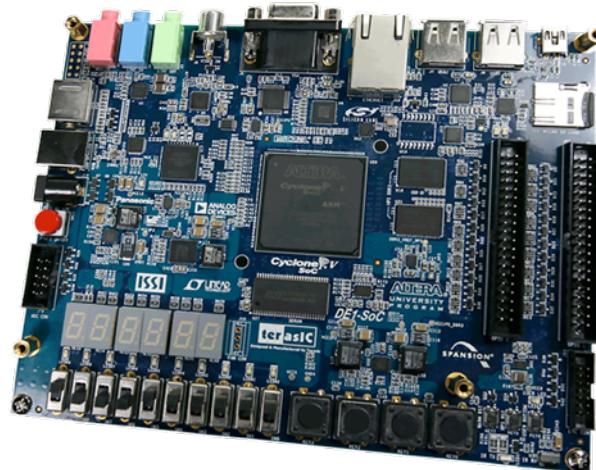
- In-person final: 40%
- Lab project: 35%
 - 3 deliverables (more on this later)
- Quizzes: 25%
 - 9 quizzes (highest 7 will be considered)
 - All online on LEARN
 - Quizzes are open book
 - You may consult your textbook, course notes, and materials posted on course webpage
 - Use of any other resource (including online services such as stackexchange.com) is prohibited
 - You may not communicate directly or indirectly with any person except course instructors (you can email course instructors if you have any questions or need any clarification)
 - You may not discuss nor disclose quiz questions with anyone

Contingency Proviso

- Course outline presents intended weights, and due dates
- As best as possible, we will keep to the outline
- We reserve the right to modify topics and/or assessments and/or weight and/or deadlines with due and fair notice
- In the event of such challenges, we will work with the Department/Faculty to find reasonable and fair solutions

Lab Project

- You will design, implement, and test
real-time executive (RTX) on DE1-SoC boards



Groups

- Groups should have 4 members
 - Never 5! 3 requires serious justification
 - Sign up in LEARN by **23:00 on Jan. 11th EST**
- Each group must have a project manager
 - Responsible for coordinating group
 - Can rotate between teammates
 - For L1, report project manager by sign-up deadline
 - For L2 and L3, report project manager by L1 and L2 submissions
- Only one split-up is allowed
 - One-week notice in writing before nearest deadline
 - All students involved lose one grace day

Milestones

Deliverable	Weight	Due Date
Group sing-up	2%	23:00 Jan 11
L1	28%	23:00 Feb 3
L2	35%	23:00 Mar 10
L3	35%	23:00 Mar 31

All times are Eastern Standard Time

Start Early!

- Time/work estimation is hard
 - Programmers are eternal optimists (it will only take two days)!
 - This is why we bug you about starting the project early
- Can a project be efficiently partitioned?
 - Partitionable task decreases in time as you add people
 - But ... what about communication?
 - Time reaches a minimum bound
 - With complex interactions, time increases!



Techniques for Partitioning Tasks

- Functional
 - Person A implements threads, Person B implements semaphores, Person C implements locks...
 - Problem: Lots of communication across APIs
 - If B changes the API, A may need to make changes
- Task
 - Person A designs, Person B writes code, Person C tests
 - May be difficult to find right balance, but can focus on each person's strengths (Theory vs systems hacker)
 - Since debugging is hard, Microsoft has two testers for each programmer

Communication

- More people means more communication
 - Changes have to be propagated to more people
 - Think about person writing code for most fundamental component of system: everyone depends on them!
- Miscommunication is common
 - “***Index starts at 0? I thought you said 1!***”
- Who makes decisions?
 - Individual decisions are fast but trouble
 - Group decisions take time
 - Centralized decisions require a big picture view (someone who can be the “system architect”)
- Often designating someone as system architect can be a good thing
 - Better not be clueless
 - Better have good people skills
 - Better let other people do work



Coordination

- Many are in different time zones ⇒ some cannot make all meetings!
 - They miss decisions and associated discussion
 - Why do we limit groups to 4 people?
 - You would never be able to schedule meetings otherwise
 - Why do we require 3 people minimum?
 - You need to experience groups to get ready for real world
- People have **different work styles**
 - Some people work in the morning, some at night
 - How do you decide when to meet or work together?
- What about project slippage?
 - Everyone busy but not talking, one is way behind, but no one will know until very end!
- Hard to add people to existing group
 - Members have already figured out how to work together



How to Make it Work?

- People are human ... get over it!
 - People will make mistakes, miss meetings, miss deadlines, etc.
 - You need to live with it and adapt
 - It is better to anticipate problems than clean up afterwards
- Document, document, document
 - Why Document?
 - Expose decisions and communicate to others
 - Easier to spot mistakes early
 - Easier to estimate progress
 - What to document?
 - Everything (but don't overwhelm people or no one will read)



Suggested Documents for You to Maintain

- Project objectives: goals, constraints, and priorities
- Specifications
 - This should be the first document generated and the last one finished
- Meeting notes
 - Document all decisions
- Schedule
 - This document is critical!
- Organizational chart
 - Who is responsible for what task?



Use Software Tools



- Source revision control software (CVS, SVN, git)
 - Easy to go back and see history
 - Figure out where and why bugs got introduced
 - Communicates changes to everyone
(use RCS's features)
- Use automated testing tools
 - Write scripts for non-interactive software
- Use E-mail and instant messaging consistently to leave history trail

Test Continuously



- Integration tests all the time, not at 8pm on due date!
 - Write dummy stubs with simple functionality
 - Schedule periodic integration tests
 - Get everyone code, build, and test ... don't wait until it is too late!
- Testing types
 - Unit tests: white-/black-box check each module in isolation
 - Daemons: subject code to exceptional cases
 - Random testing: subject code to random timing changes
- Test early, test later, test again
 - What if something changes in some other part of code?

Late Submissions

- 3 grace days (including weekends) without penalty
- 15% per day late submission penalty afterwards
 - 1-hour-late submission = 15-hour-late submission
- Late submissions are not accepted after three days

Collaboration Policy

- Explaining concepts to someone in another group
 - Discussing algorithms/testing strategies with other groups
 - Helping debug someone else's code (in another group)
 - Searching online for generic algorithms (e.g., hash table)
-
- Sharing code or test cases with another group
 - Open-sourcing code (e.g., on GitHub) even after this term
 - Copying OR reading another group's code or test cases
 - Copying OR reading online code or test cases from prior years
-
- Zero tolerance policy for plagiarism
 - We use [Moss](#) and follow [UW Policy 71](#) for any single incident



Seeking Help

- Lab Q&A on Piazza discussion forum
 - Looking for group partners
 - Lab/Project administration
 - ARM DS Q&A
 - Project Q&A
 - Target response time: one business day
 - **Do not wait till the last minute to ask questions**
- Individual emails
 - Only for questions containing confidential information
- Office hours
- Appointment

Important Near-term Task

Sign up for project groups on LEARN by
23:00 on January 11th, 2022 EST

WORK HARD



RELAX HARDER

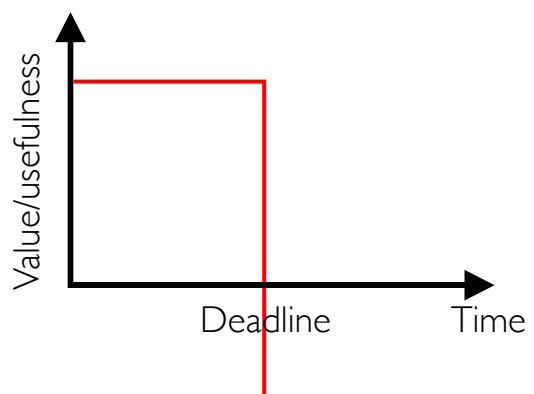
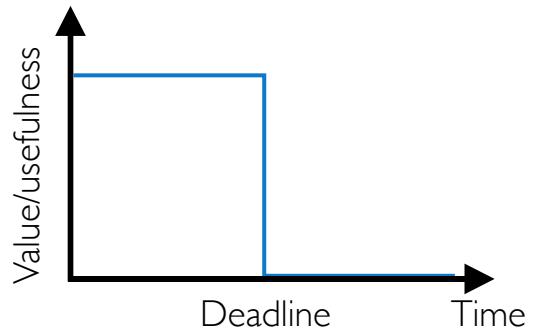
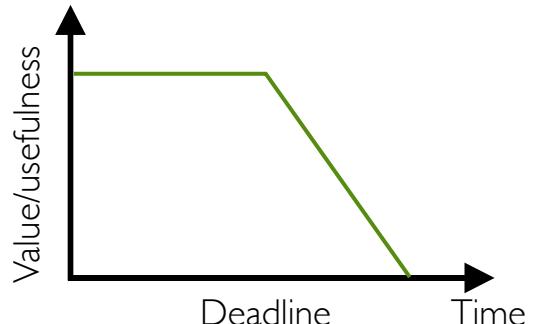
makeameme.org

What is a Real-time System?

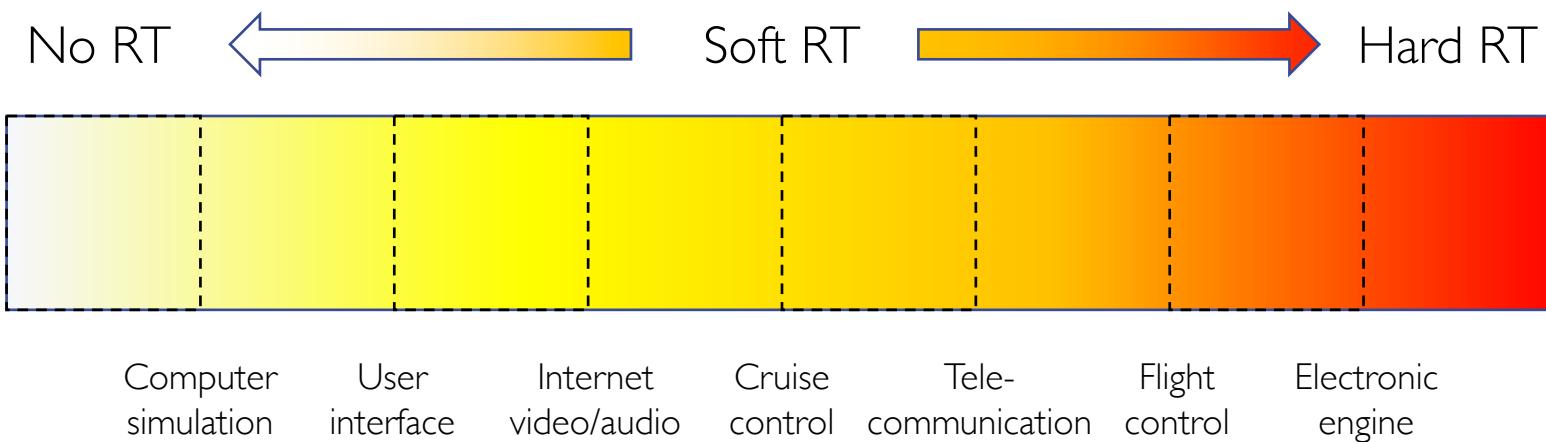
- Systems whose correctness depends on both their **temporal** aspects as well as their **functional** aspects
- Temporal aspects are typically specified through **deadline**
 - Results must be produced before deadline
 - Correct result produced too late is considered failure
- Real-time systems are typically embedded control systems
 - E.g., safety-critical systems
- OS must be carefully designed to support temporal properties

Types of Real-time Systems

- **Soft**: must try to meet all deadlines
 - System does not fail if a few deadlines are missed
- **Firm**: result has no use outside deadline window
 - Tasks that fail are discarded
- **Hard**: must always meet all deadlines
 - System fails if deadline window is missed



Real-time Spectrum

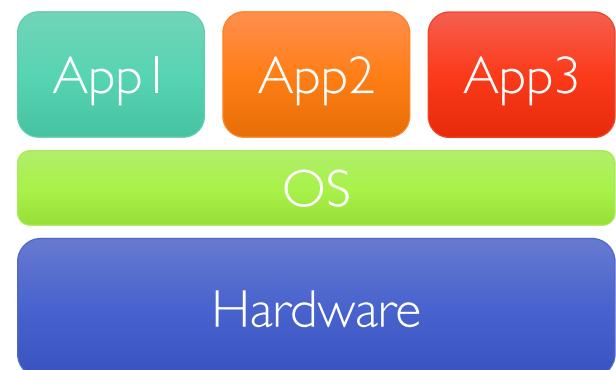


What is an Operating System?

- No universally accepted definition
- “Everything vendors ship when you order OS” is good approximation, but varies wildly
- “The one program running at all times on computer” is kernel
 - Everything else is either system program (ships with OS) or application program

What is an Operating System? (cont.)

- Special layer of software that provides applications access to hardware resources
 - **Abstract** view of complex hardware devices
 - **Protected** access to shared resources
 - Security and authentication
 - Communication amongst logical entities

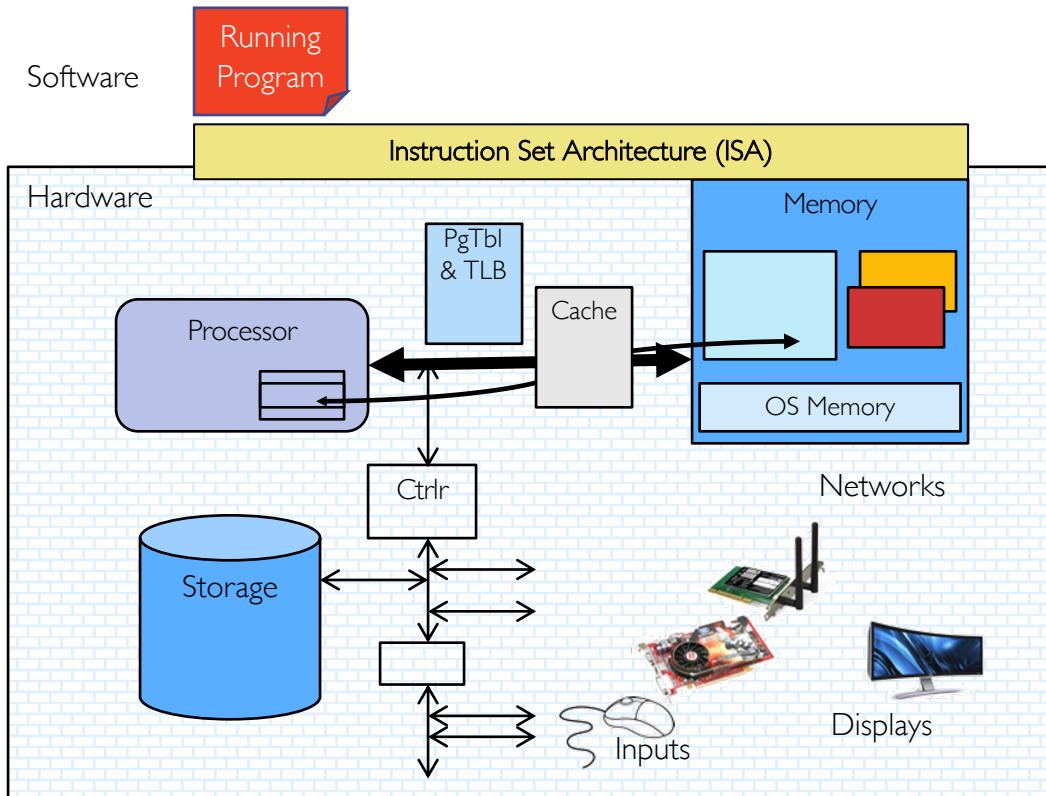


What is an Operating System? (cont.)

- Illusionist
 - Provide clean, easy-to-use abstractions of physical resources
 - Infinite memory, dedicated machine
 - Higher level objects: files, users, messages
 - Masking limitations, virtualization

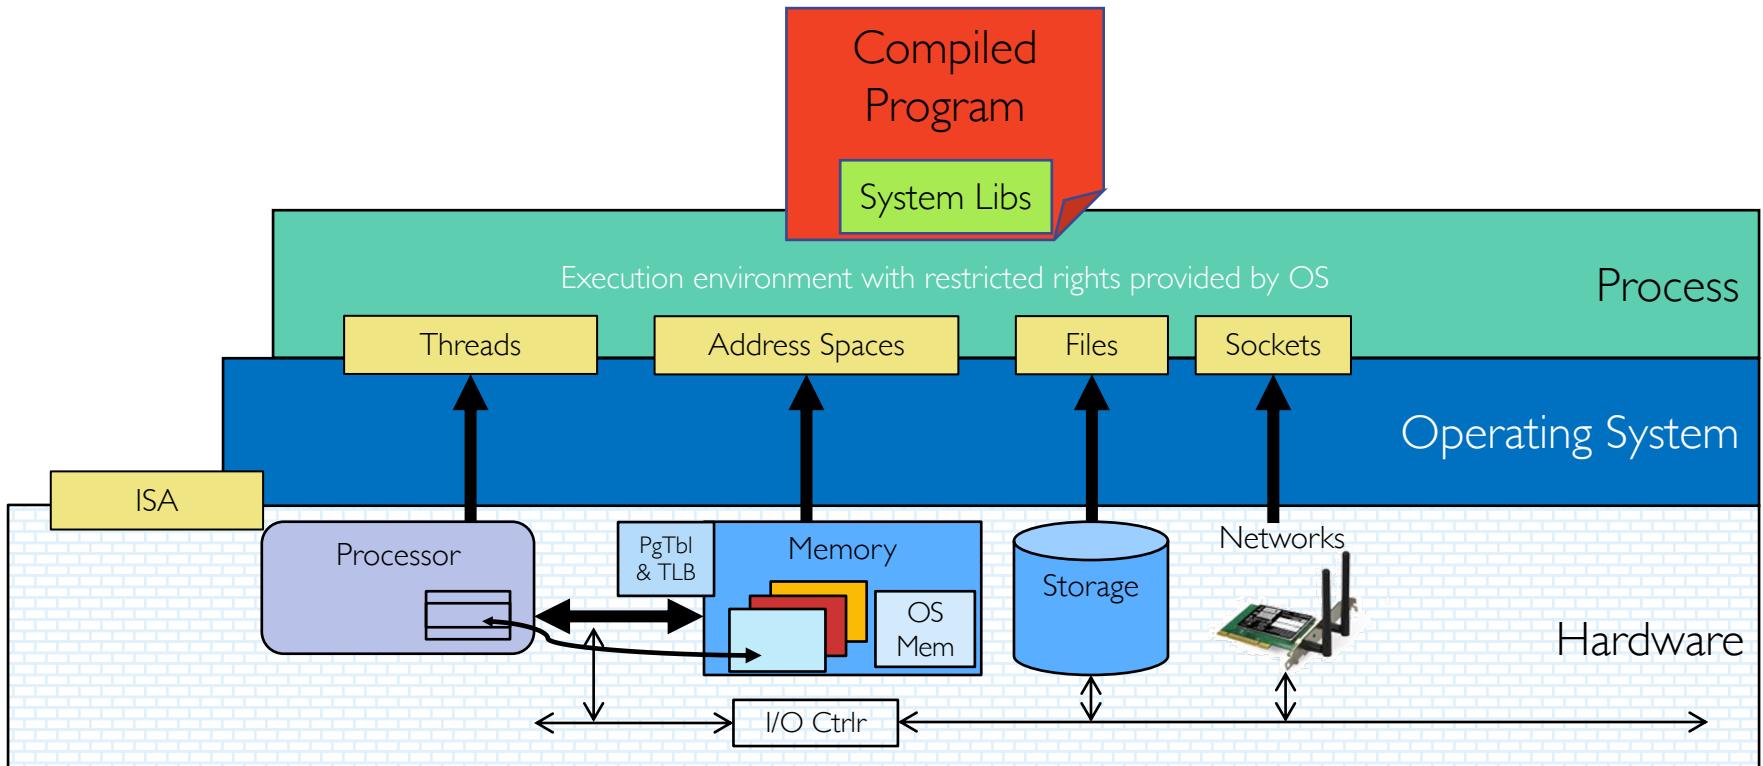


Hardware/Software Interface



- ECE 222 and ECE 320: Machine structures (and C)
- OS *abstracts* these hardware details from the application

OS Basics: Virtualizing Hardware

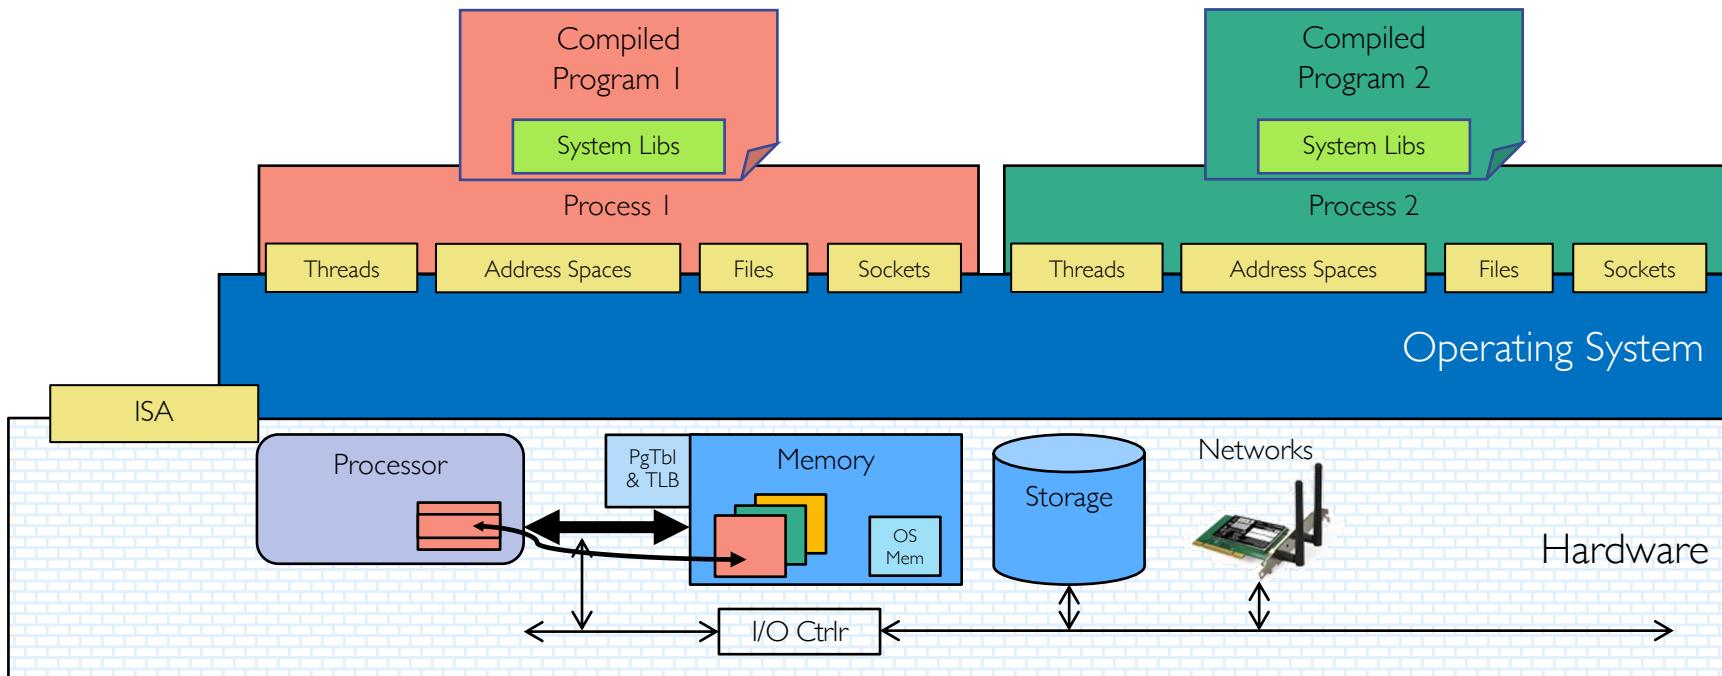


What is an Operating System? (cont.)

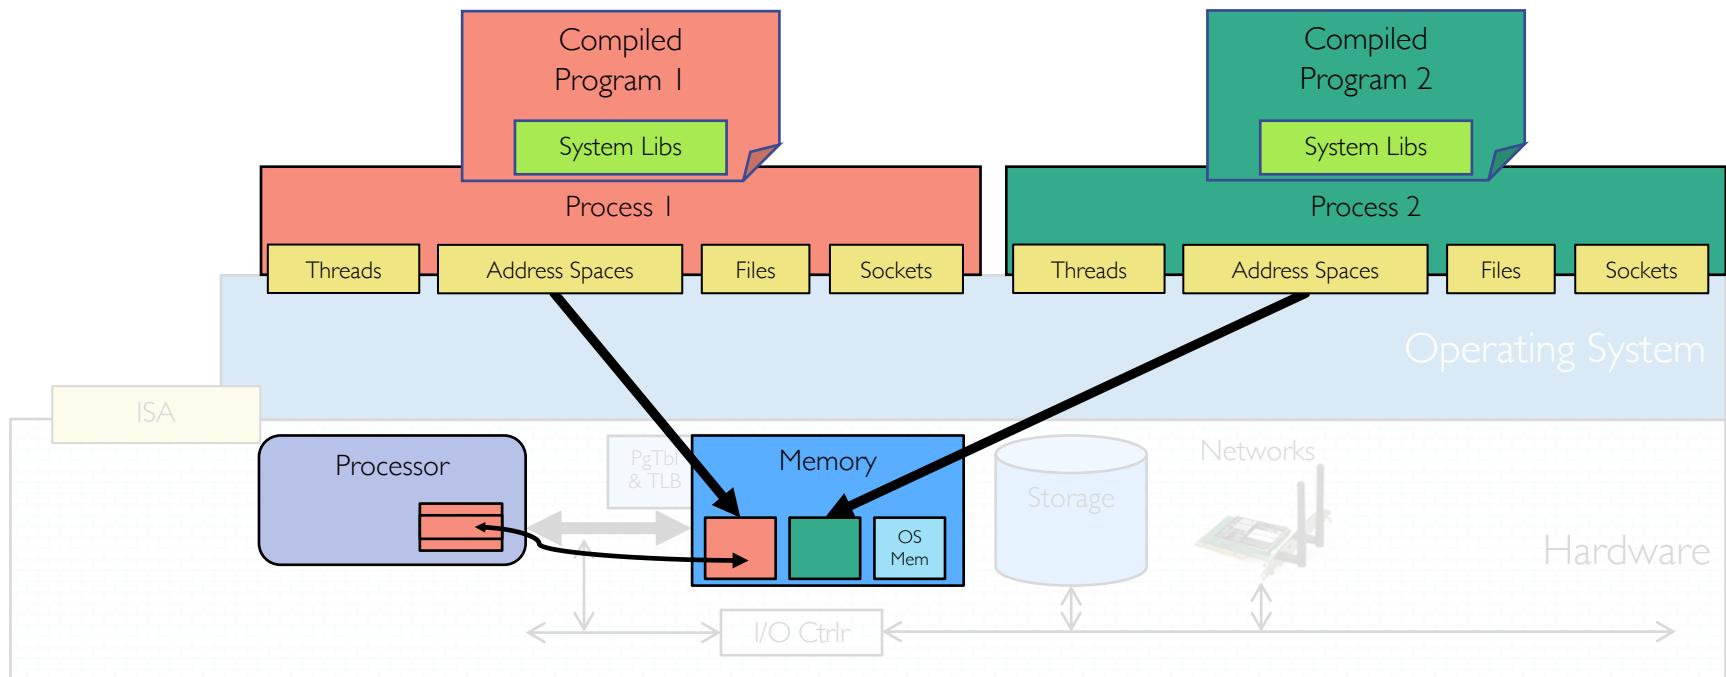
- Illusionist
 - Provide clean, easy-to-use abstractions of physical resources
 - Infinite memory, dedicated machine
 - Higher level objects: files, users, messages
 - Masking limitations, virtualization
- Referee
 - Provide protection, isolation, and sharing of resources
 - Resource allocation and communication



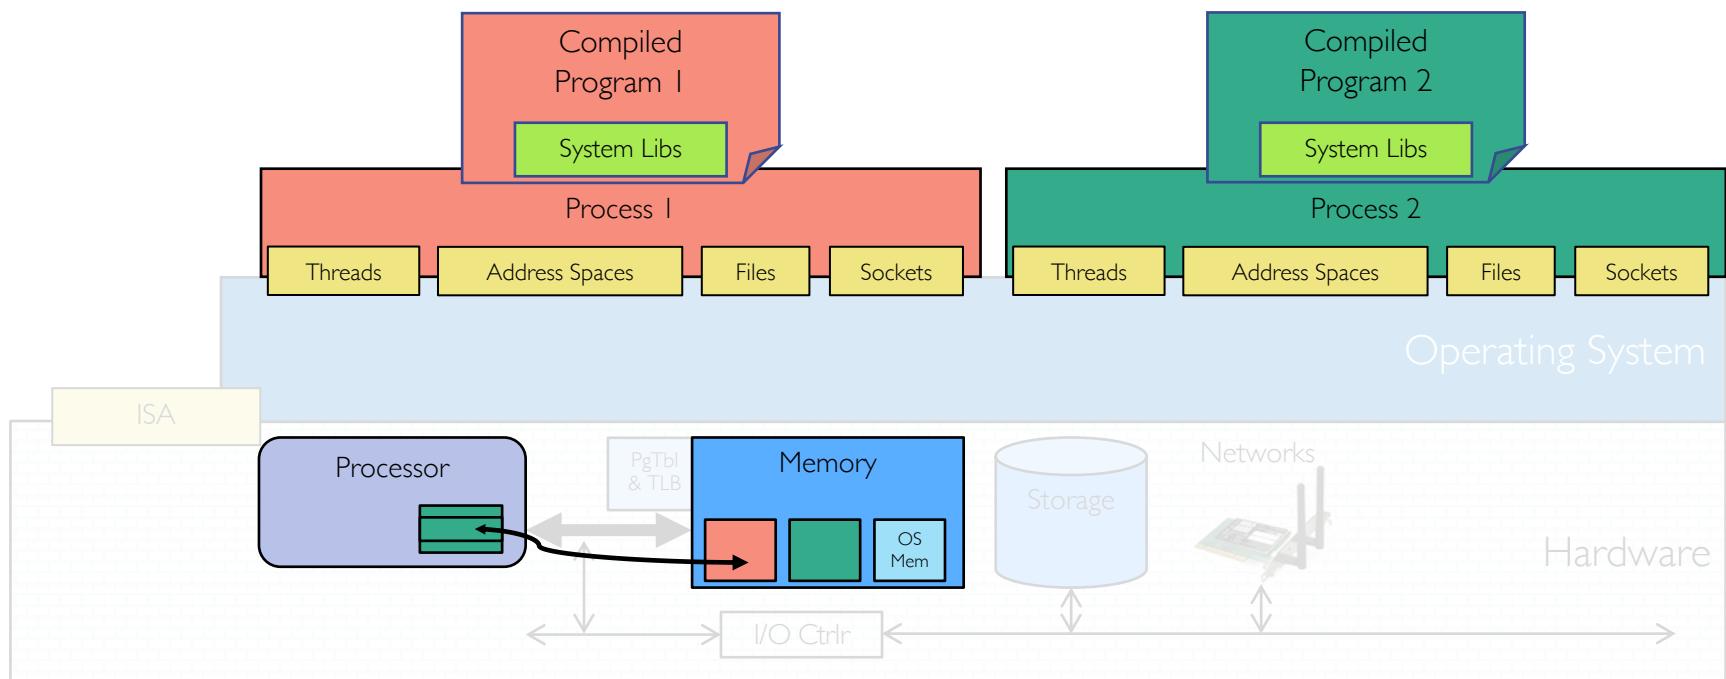
OS Basics: Switching Processes



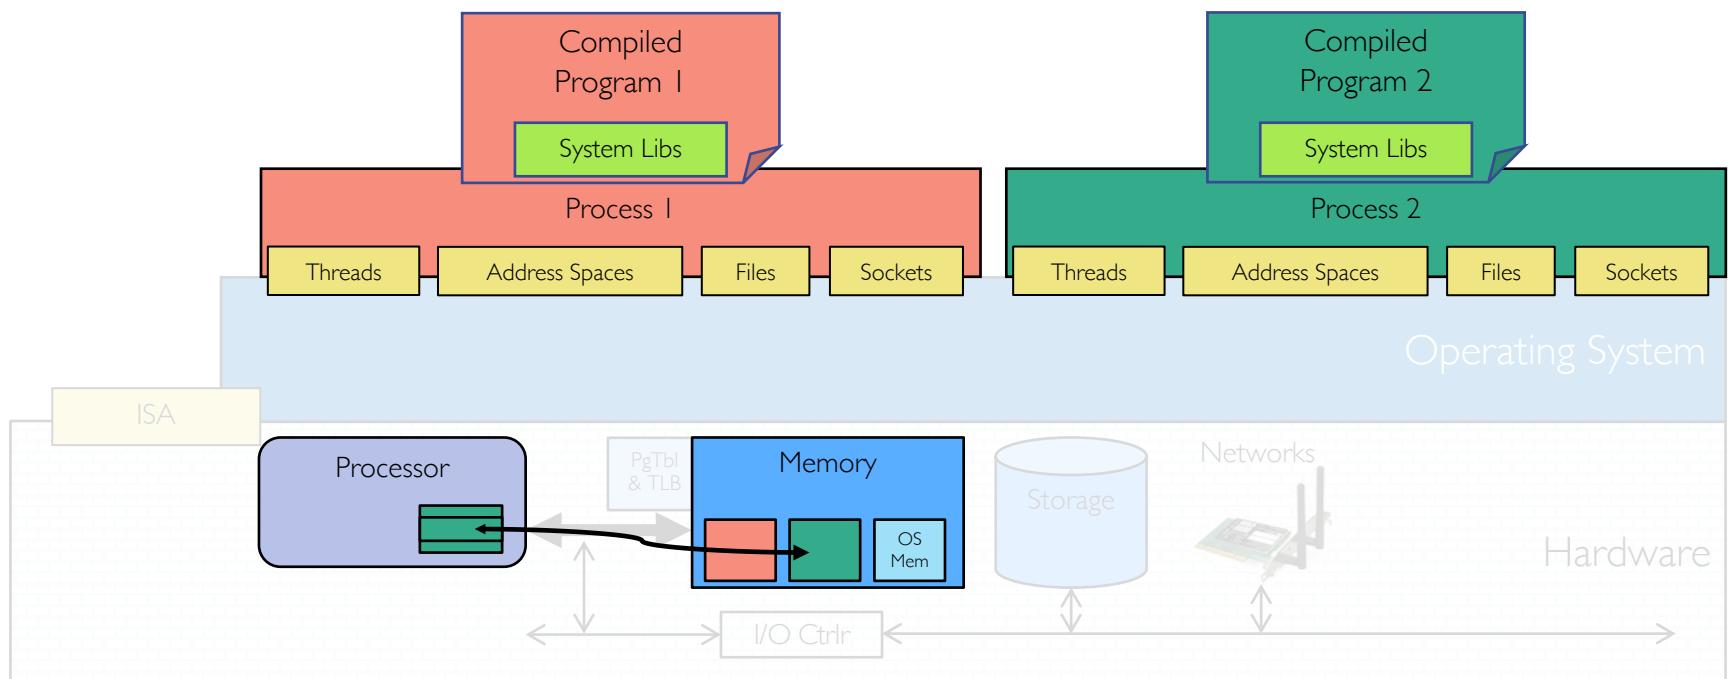
OS Basics: Switching Processes



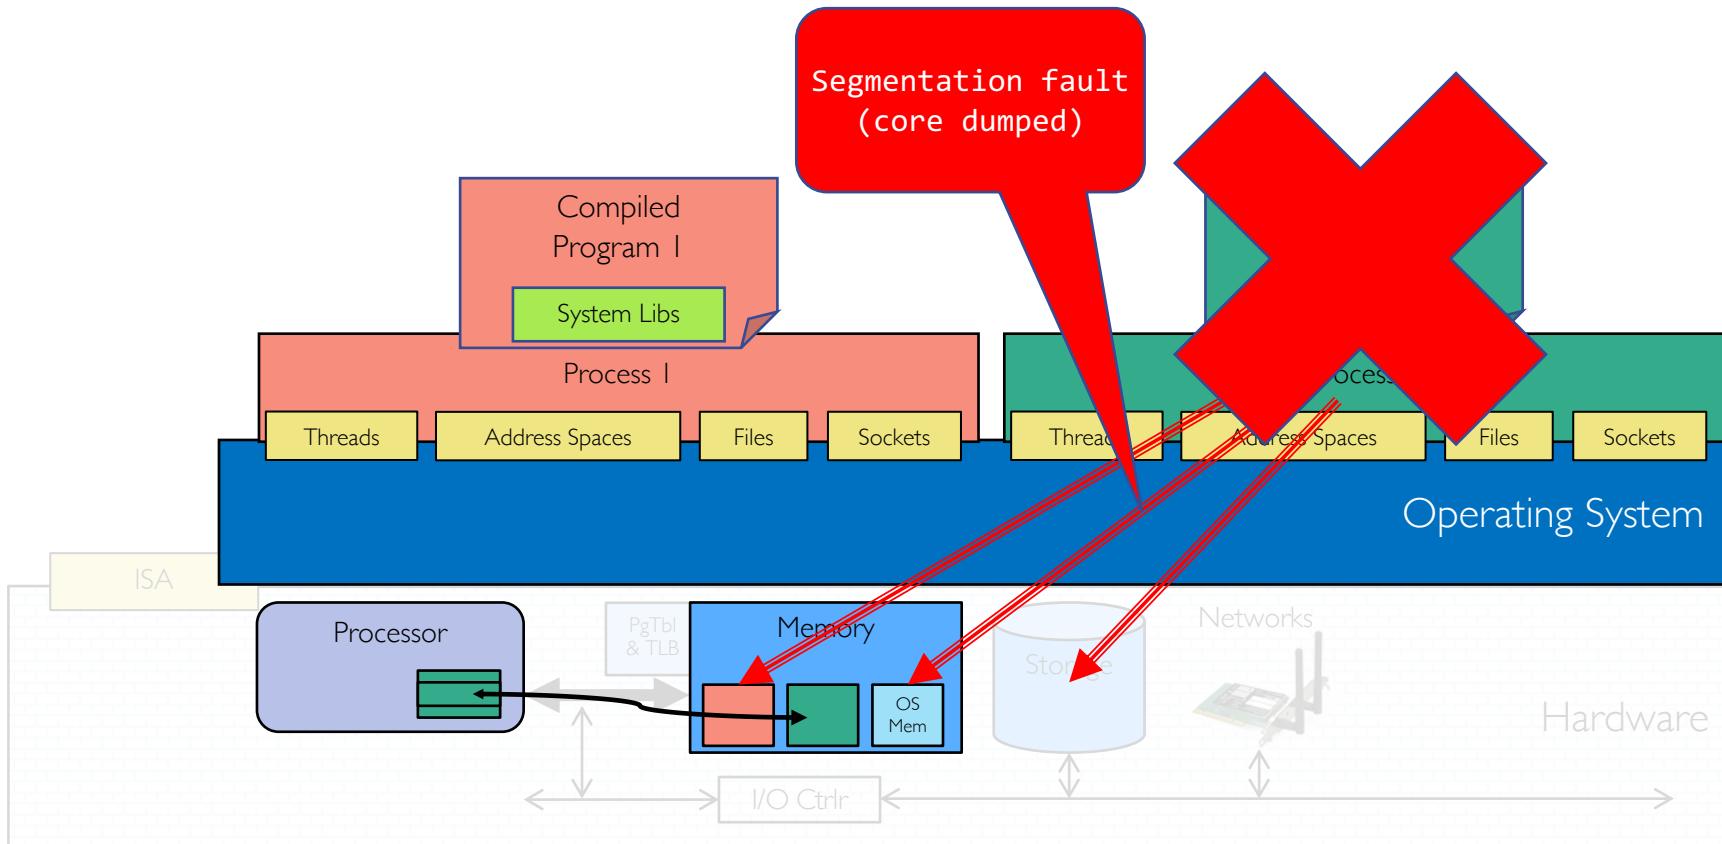
OS Basics: Switching Processes (cont.)



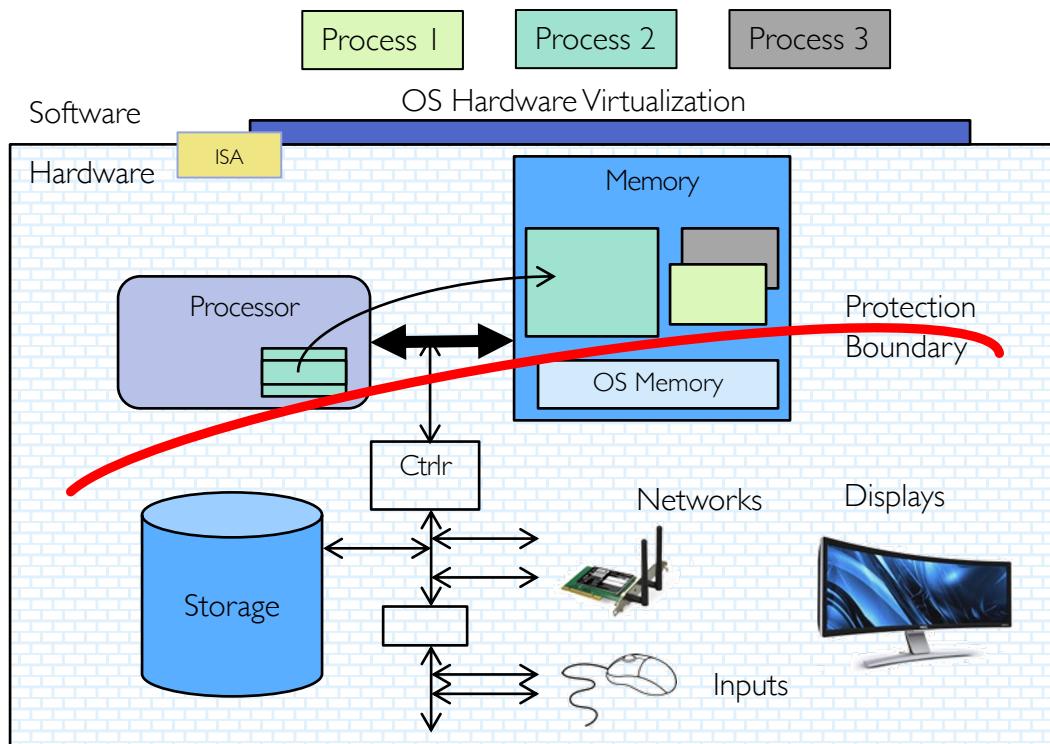
OS Basics: Switching Processes (cont.)



OS Basics: Protection



OS Basics: Protection (cont.)



- OS isolates processes from each other
- OS isolates itself from other processes
- ... even though they run on the same HW!

What is an Operating System? (cont.)

- Illusionist



- Provide clean, easy-to-use abstractions of physical resources
 - Infinite memory, dedicated machine
 - Higher level objects: files, users, messages
 - Masking limitations, virtualization

- Referee



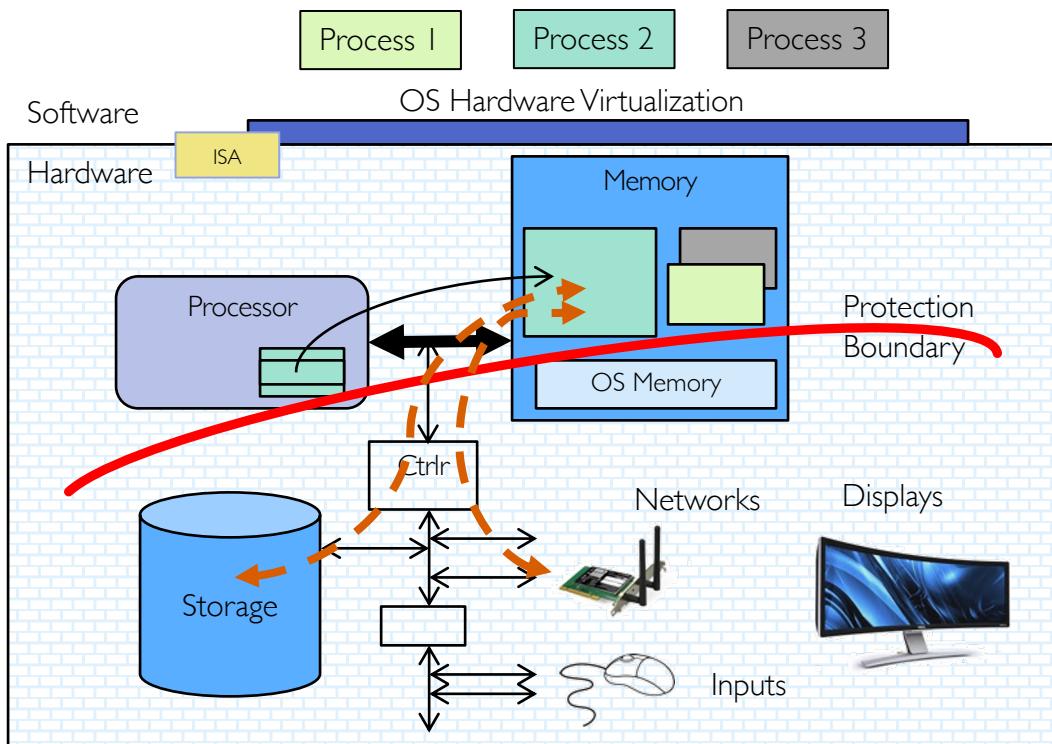
- Provide protection, isolation, and sharing of resources
 - Resource allocation and communication

- Glue



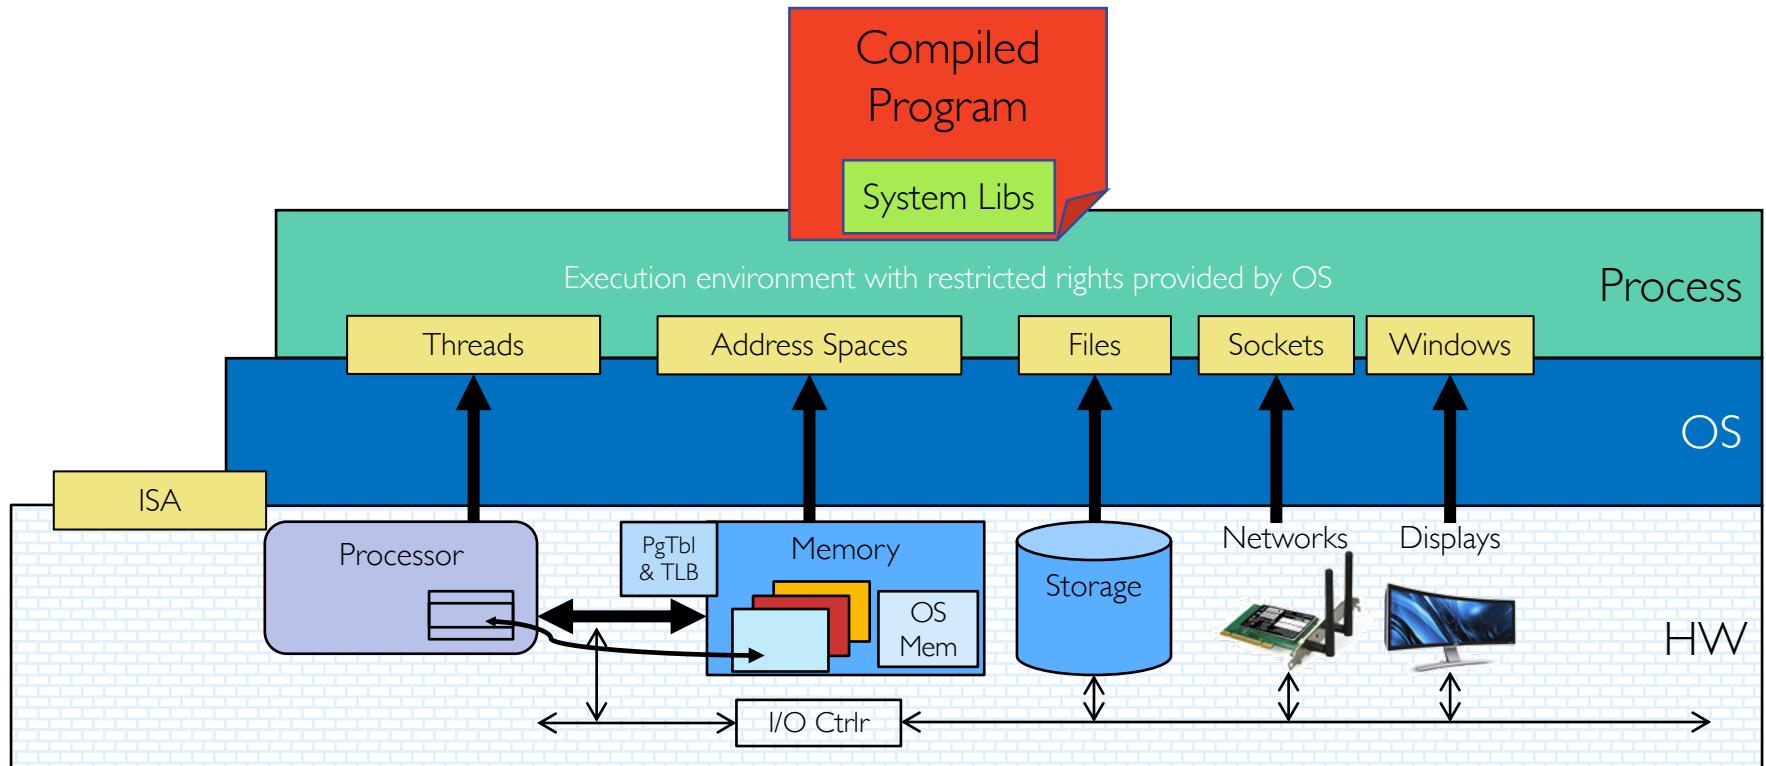
- Provide common services
 - Storage, window system, networking, sharing, authorization
 - Look and feel

OS Basics: I/O

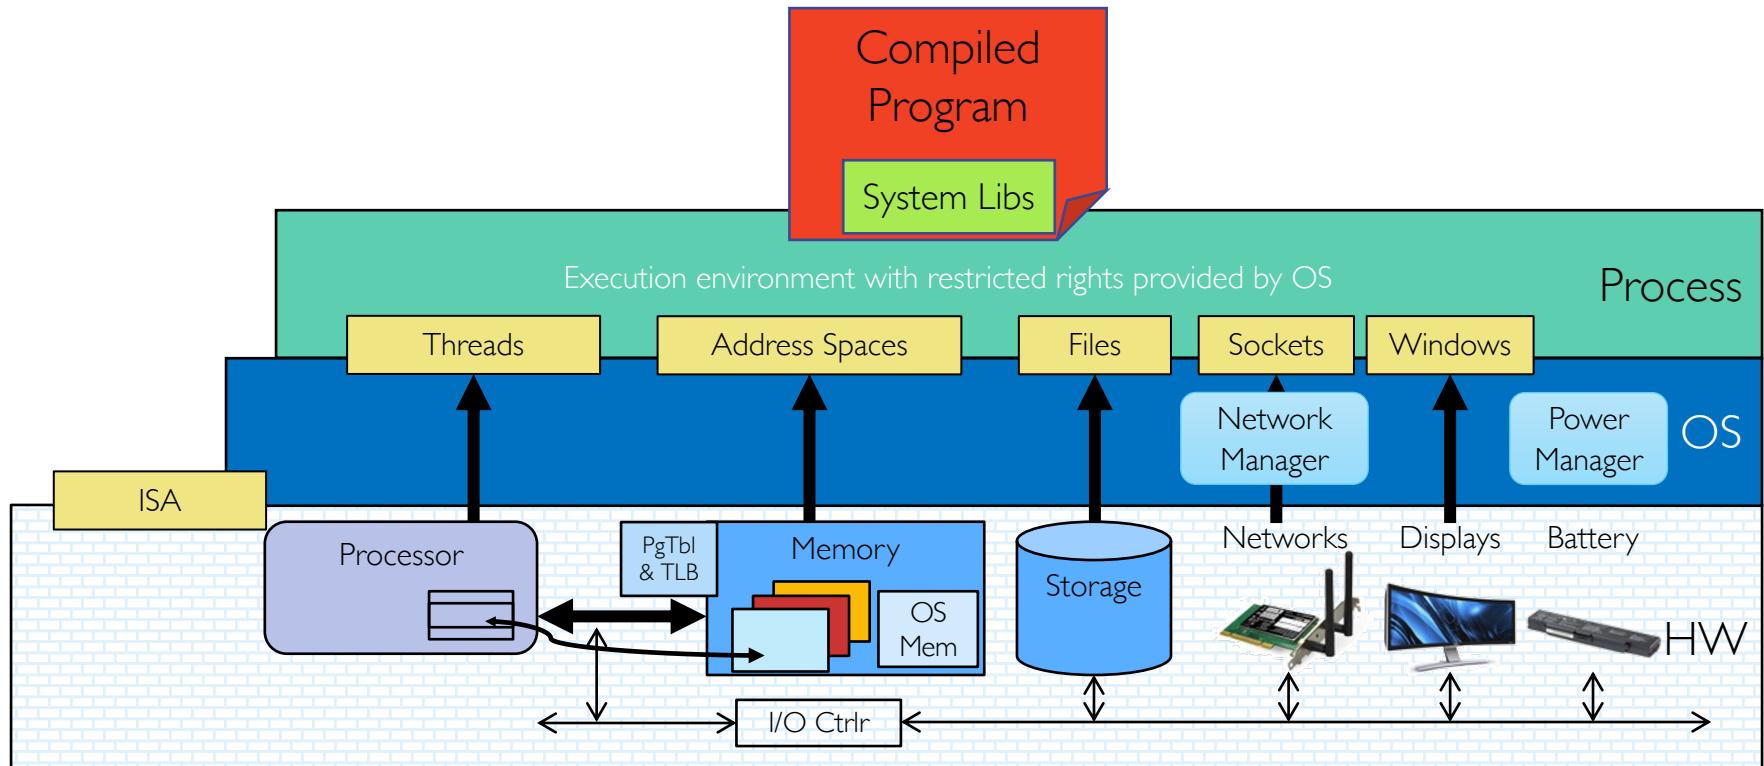


- OS provides common services in the form of I/O

OS Basics: Look and Feel



OS Basics: Background Management

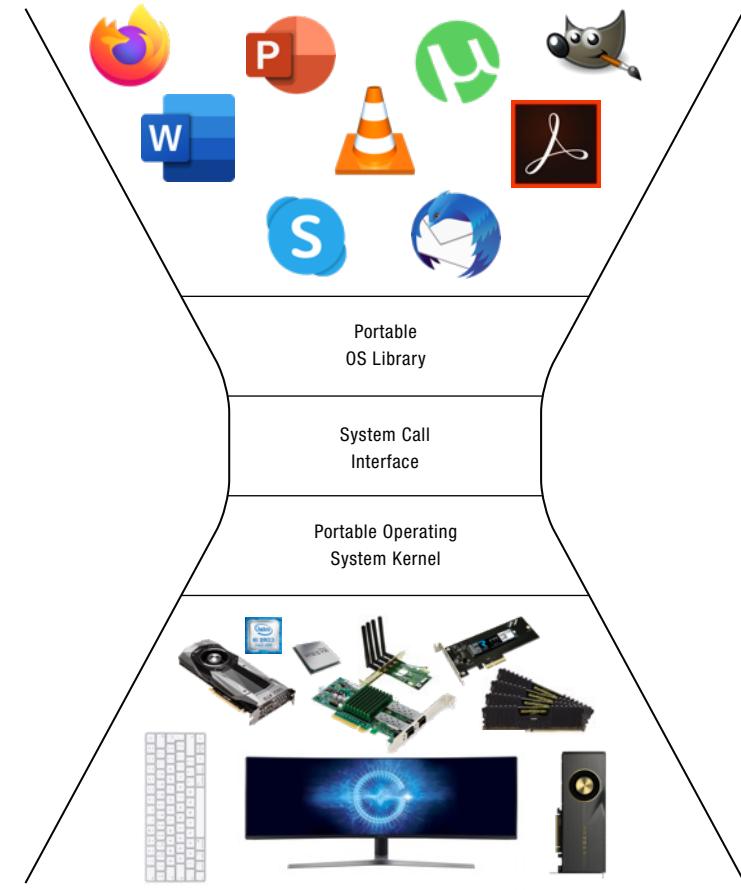


OS Basics: Hardware Support

- OS bottom line is to support applications!
 - OS itself is **incidental**
 - Ideally, OS should have very low performance overhead over raw hardware
- OS relies on HW support to provide abstractions **efficiently**
 - Dual-mode operation, interrupts, traps, precise exceptions, memory management unit, translation lookaside buffer, etc.
- HW support and OS design continue to co-evolve...
 - ... as hardware performance improves (e.g., faster storage/network), ...
 - ... and application requirements change
 - What we study in this class is result of decades of co-evolution!

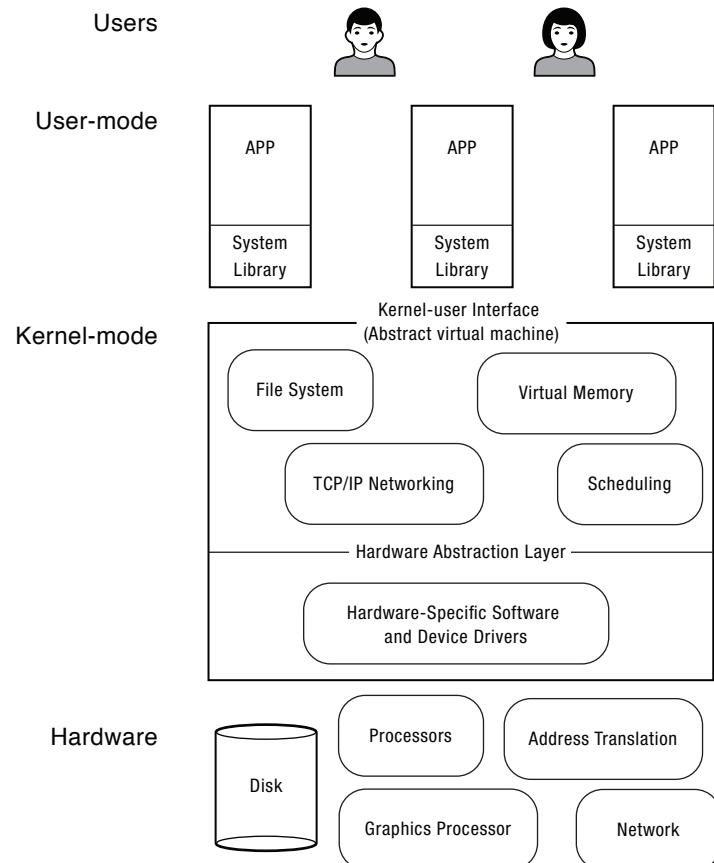
What Do Operating Systems Do?

- Provide abstractions to applications
 - File systems
 - Processes, threads
 - Virtual memory
 - Naming system, ...
- Manage diverse resources
 - Memory, CPU, storage, ...
- Achieves above by implementing specific algorithms and techniques
 - Scheduling
 - Concurrency
 - Transactions
 - Security, ...



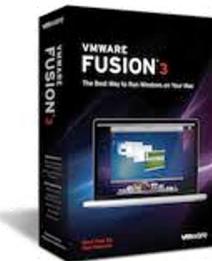
What Do Operating Systems Do? (cont.)

- Manage hardware resources for users and applications
- Convert what hardware gives into something that application programmers want
- For any OS component, begin by asking two questions
 - What is hardware interface? (physical reality)
 - What is application interface? (virtual machine)



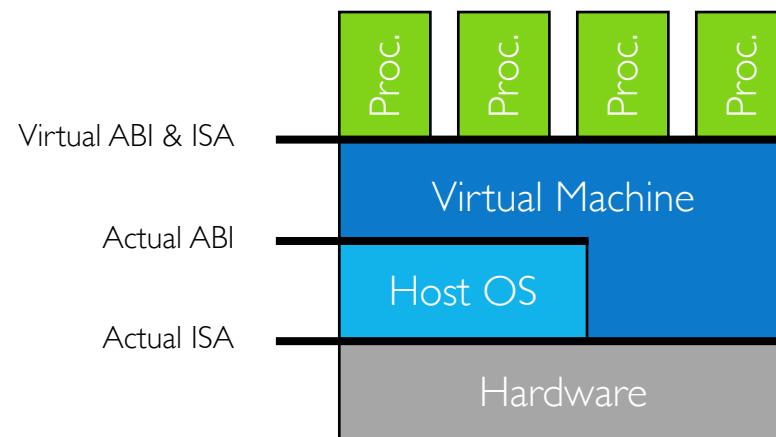
Virtual Machines (VMs)

- Software that emulates physical machine
 - Gives programs **illusion** that they run on physical machine
 - Provides platform that is independent of actual underlying hardware
 - Makes it look like hardware has features programs want
- Two types of virtual machines
 - Process VM: supports execution of single program (e.g., Java)
 - System VM: supports execution of entire OS (e.g., VMWare Fusion, Virtual box, Parallels Desktop, Xen)



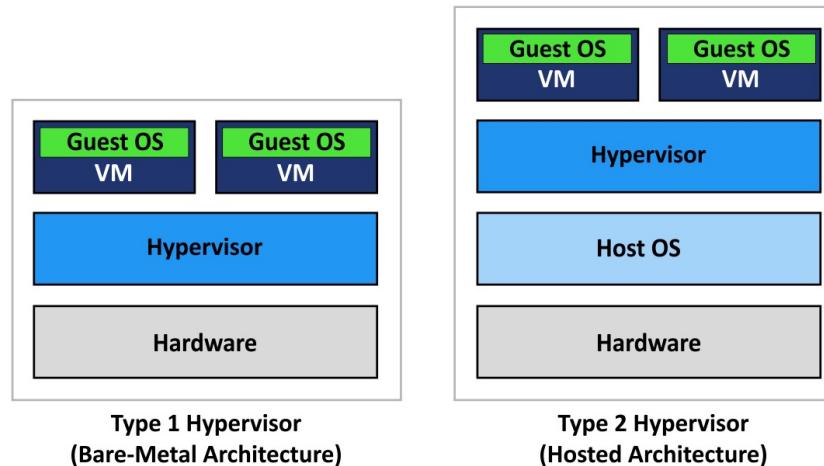
Process VMs

- Runs processes
 - Abstracts underlying OS and hardware
 - Provides platform-independent environment
 - E.g., Java virtual machine, .NET framework



System Virtual Machines: Layers of OSes

- Runs OSes
 - Useful for OS development and testing programs on other OSes
- Hypervisors create and run virtual machines
- **Type-I** hypervisors allocate HW to VMs in addition to managing them
 - E.g., Xen, VMWare ESXi
- **Type-II** hypervisors rely on host OS for HW management
 - E.g., Virtual Box, VMWare Workstation, KVM

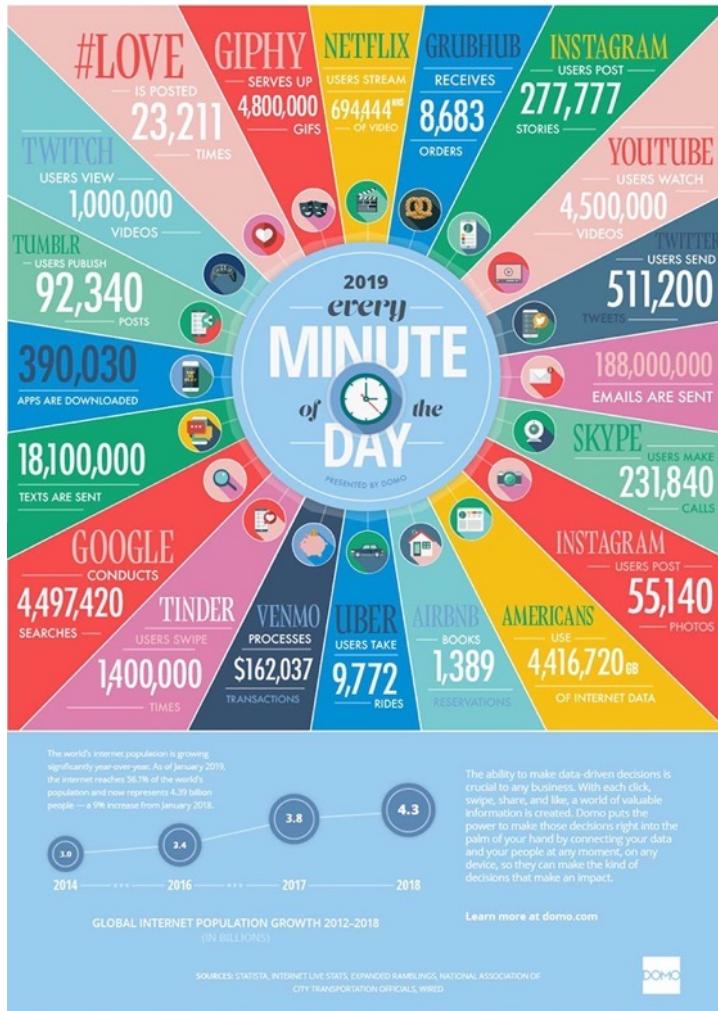


Containers: Low-weight Alternatives to Full-system Virtualization

- Provide OS virtualization above single shared kernel
 - Do not provide *full-machine* virtualization
 - Each VM has illusion of running on isolated machine
 - Each container has illusion of running on isolated OS
- Use OS constructs to provide *sand boxes* for execution
 - E.g., Linux cgroups, namespaces, etc.
- Can run on bare metal OS, or atop of OS running in VM
- **OS containers:** multiple applications run in same container
 - E.g., LXC, OpenVZ, FreeBSD Jail
- **Application containers:** each application has its own container
 - E.g., Docker, rkt



What Makes Operating Systems so Exciting and Challenging?

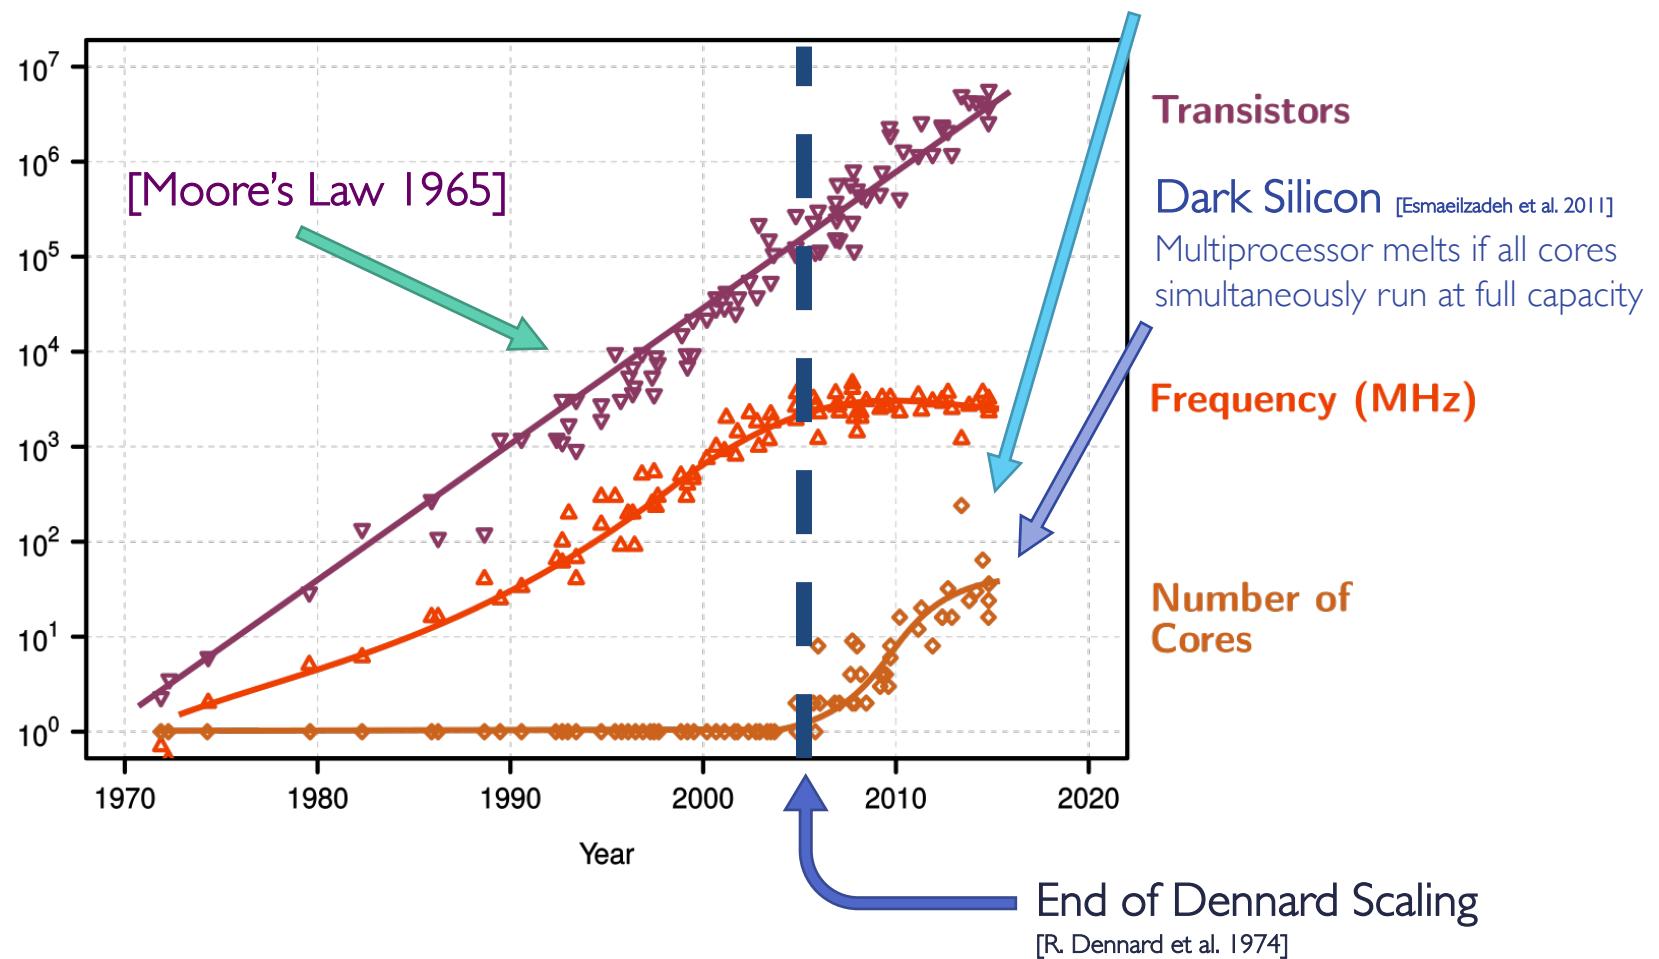


Operating systems are at the heart of it all ...

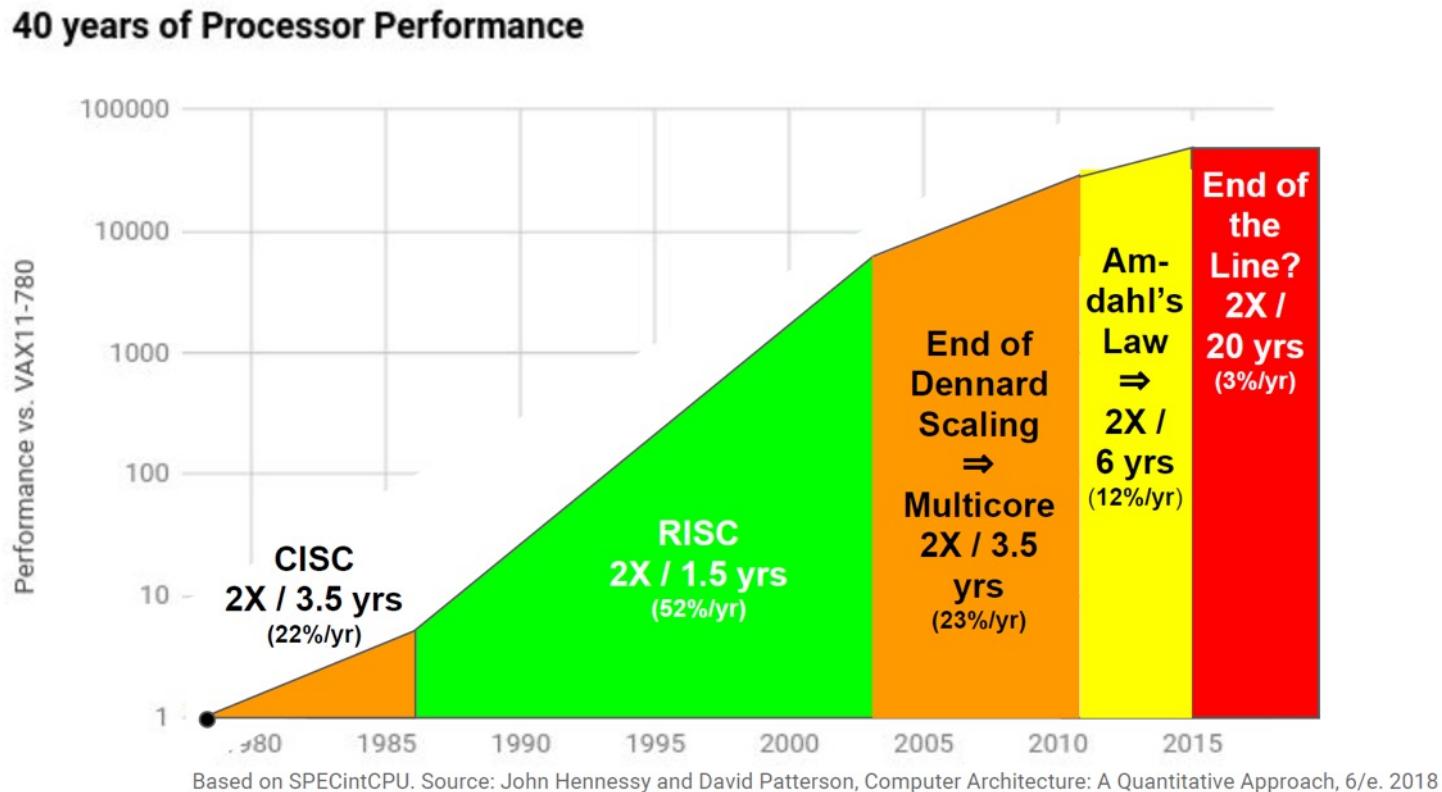
- Challenges
 - Keeping up with evolving HW
 - Managing ever-growing complexity of SW

Technology Trends

How do we program these?
Parallelism must be exploited at all levels

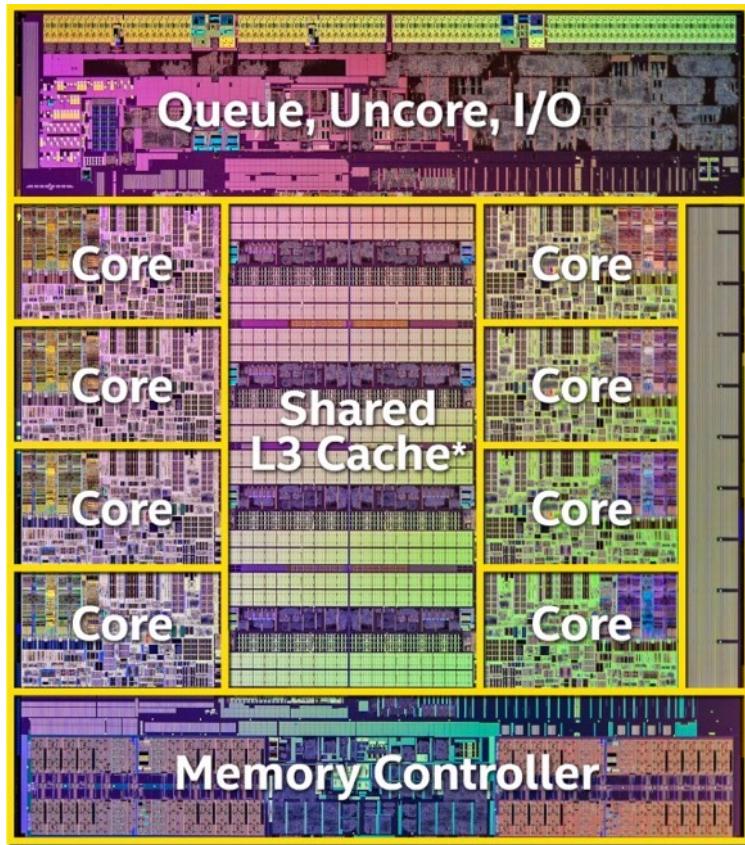


End of Growth of Single Program Speed



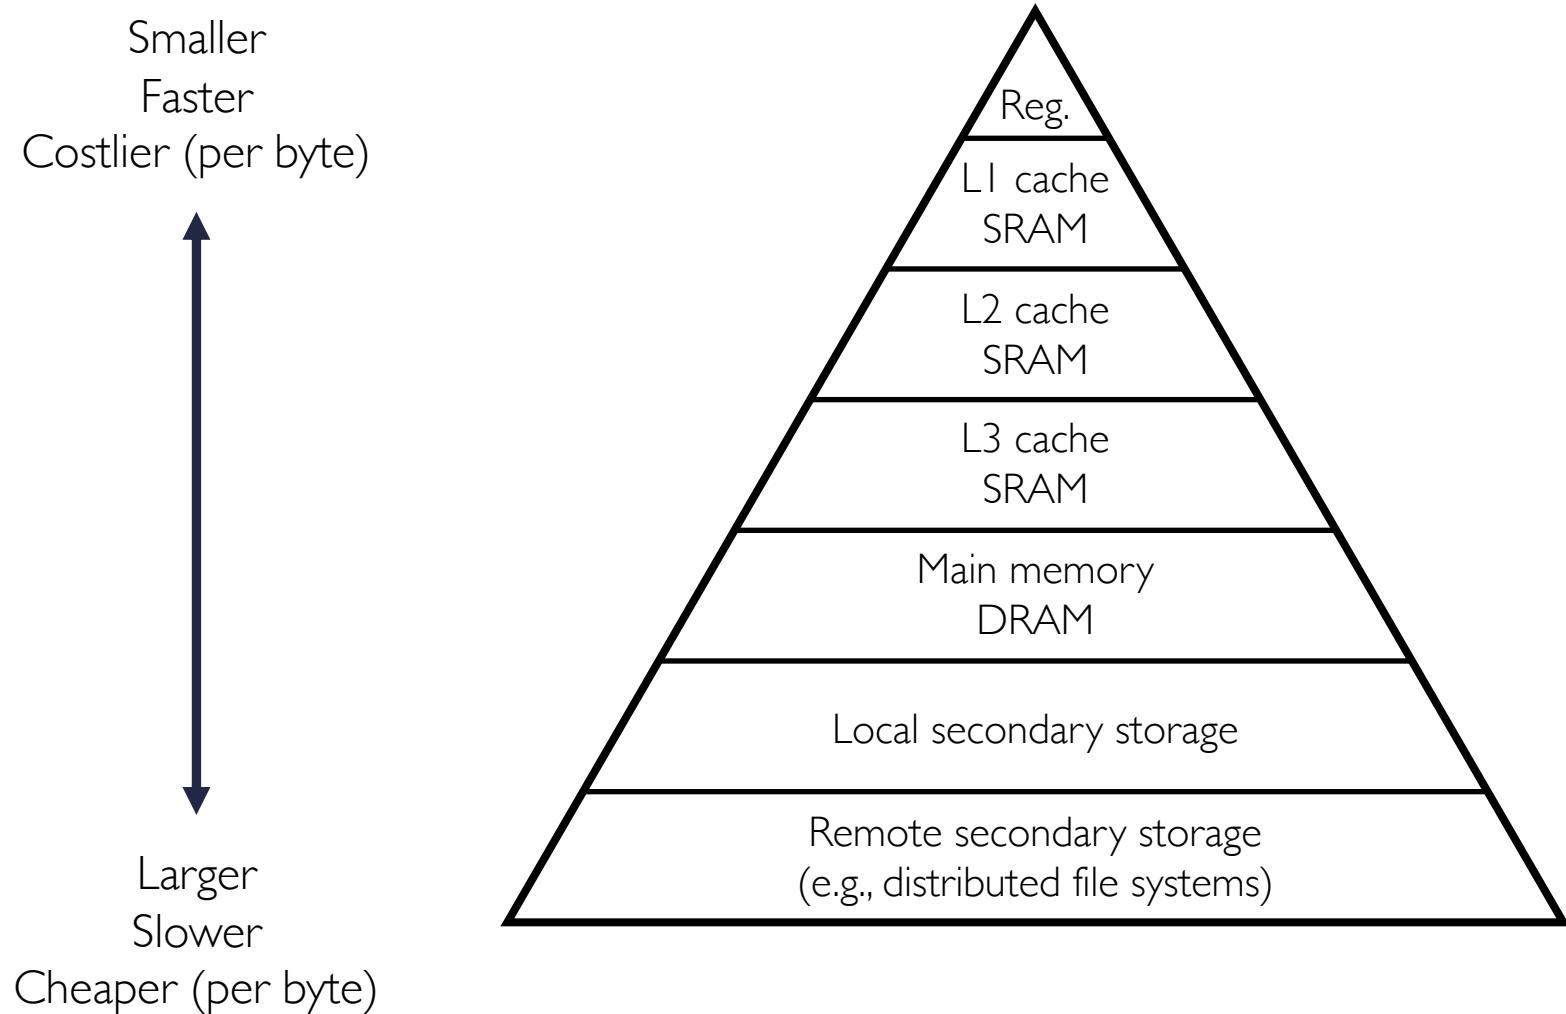
Modern Processors

Intel Haswell E



- Intel Xeon Platinum 9282
 - 14nm processor
 - 56 cores, 112 threads
 - 1.75MB data and ins. L1 cache
 - 56MB L2 cache
 - 77MB shared L3 cache
 - 8B transistors
- AMD EPYC 7H12
 - 7nm processor
 - 64 cores, 128 threads
 - 2MB data and ins. L1 cache
 - 32MB L2 cache
 - 256MB shared L3 cache
 - 4.8B transistors

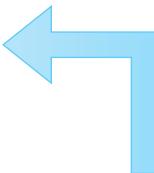
Memory Hierarchy



Numbers Everyone Should Know

[Jeff Dean, 2009]

L1 cache reference	0.5 ns
Branch mispredict	5 ns
L2 cache reference	7 ns
Mutex lock/unlock	25 ns
Main memory reference	100 ns
Compress 1K bytes with Zippy	3,000 ns
Send 2K bytes over 1 Gbps network	20,000 ns
Read 1 MB sequentially from memory	250,000 ns
Round trip within same datacenter	500,000 ns
Disk seek	10,000,000 ns
Read 1 MB sequentially from disk	20,000,000 ns
Send packet CA->Netherlands->CA	150,000,000 ns



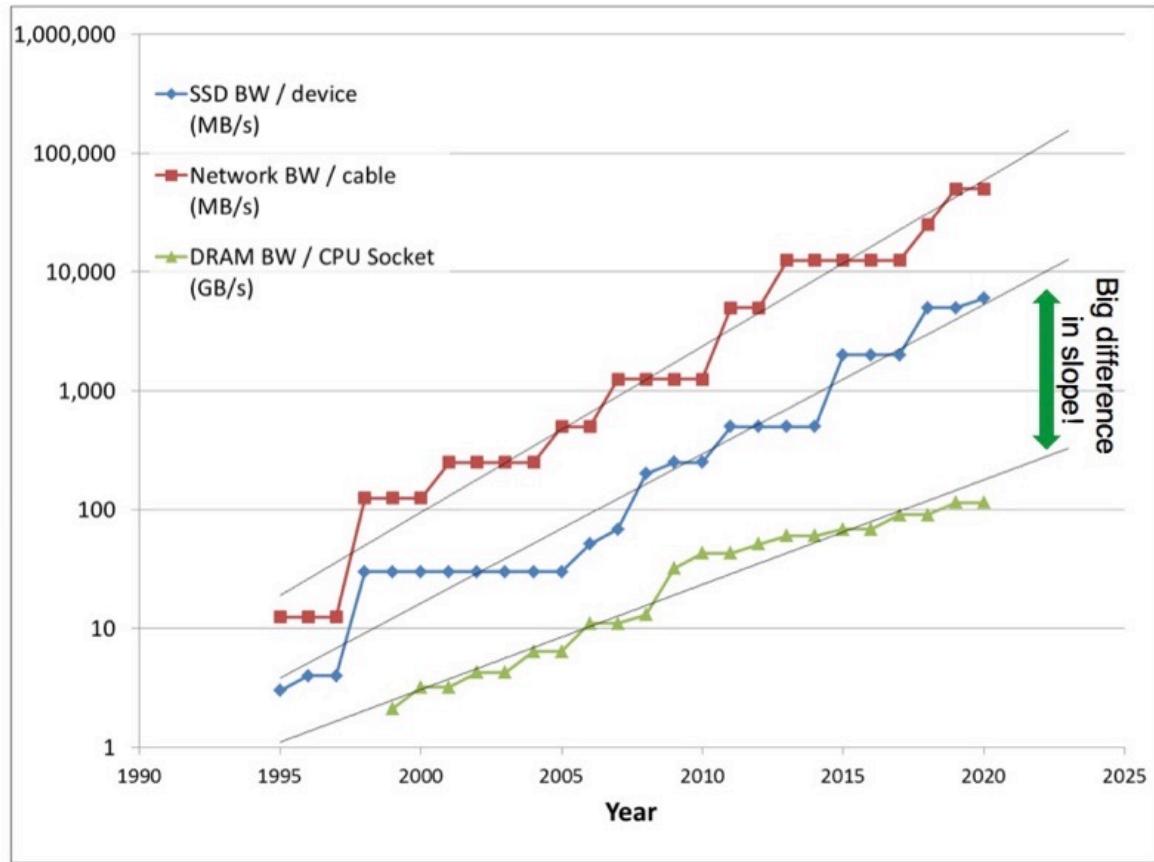
Key stroke ~100 ms

Network, IO, and Memory Bandwidth Trends

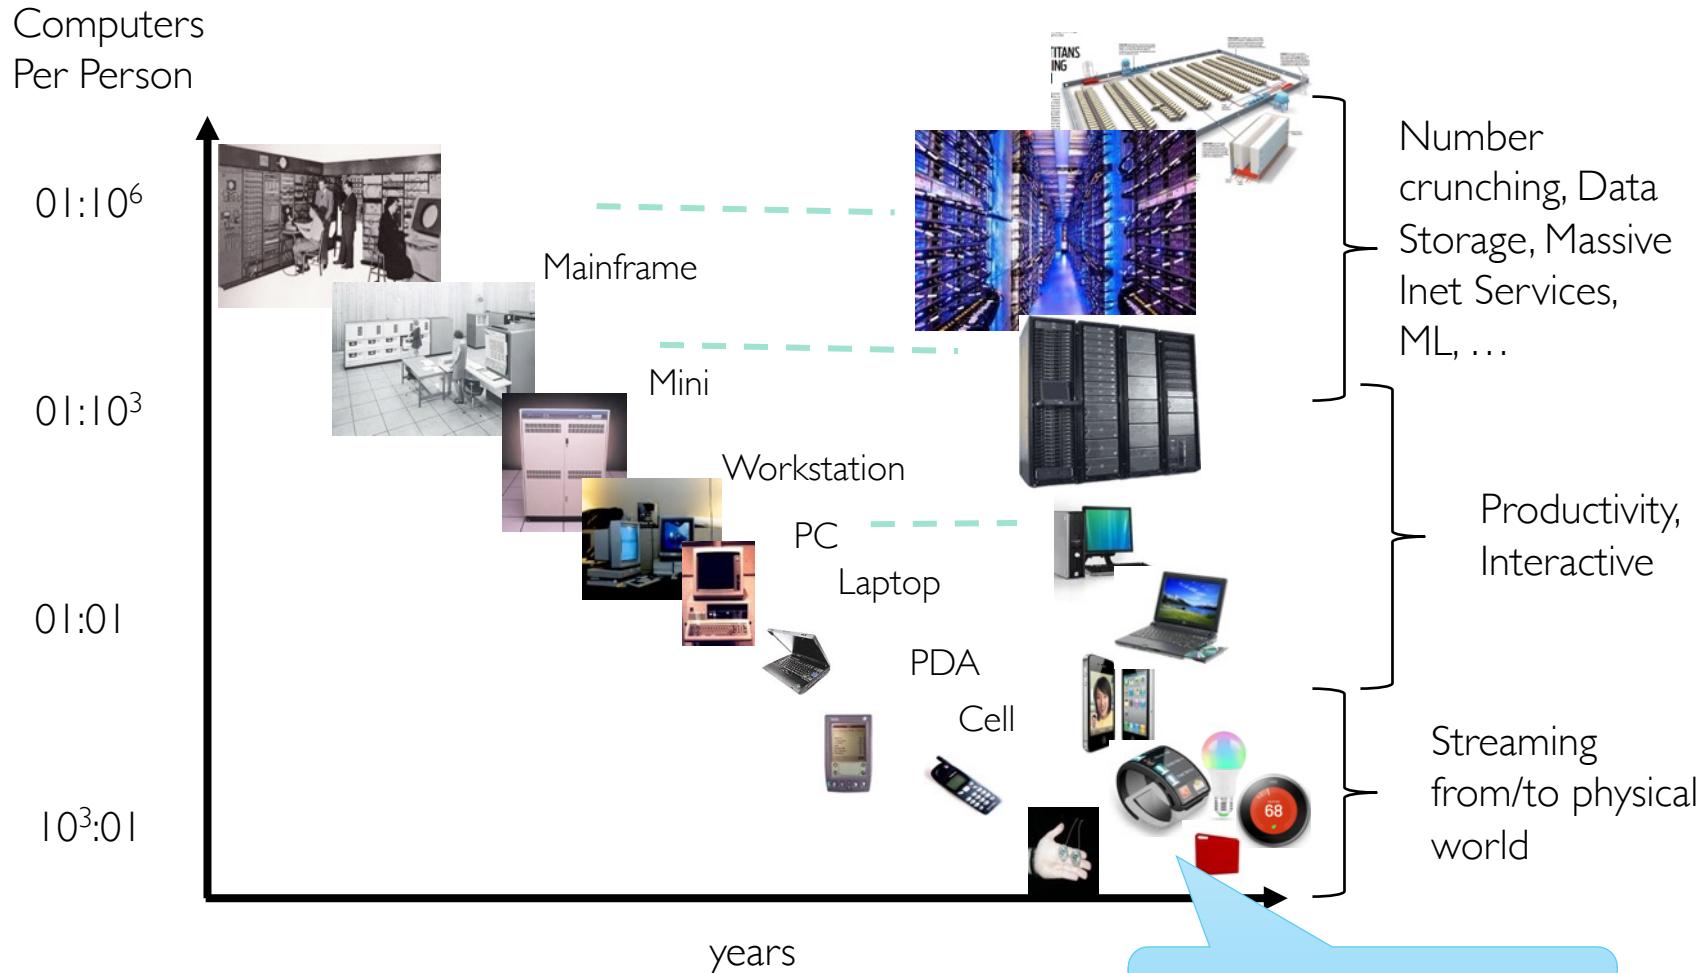
Network, Storage, & DRAM trends

Log scale

- Use DRAM Bandwidth as a proxy for CPU throughput
- Reasonable approximation for DMA and poor cache performance workloads (e.g. Storage)



People to Computer Ratio Trend



Bell's Law: new computer class per 10 years

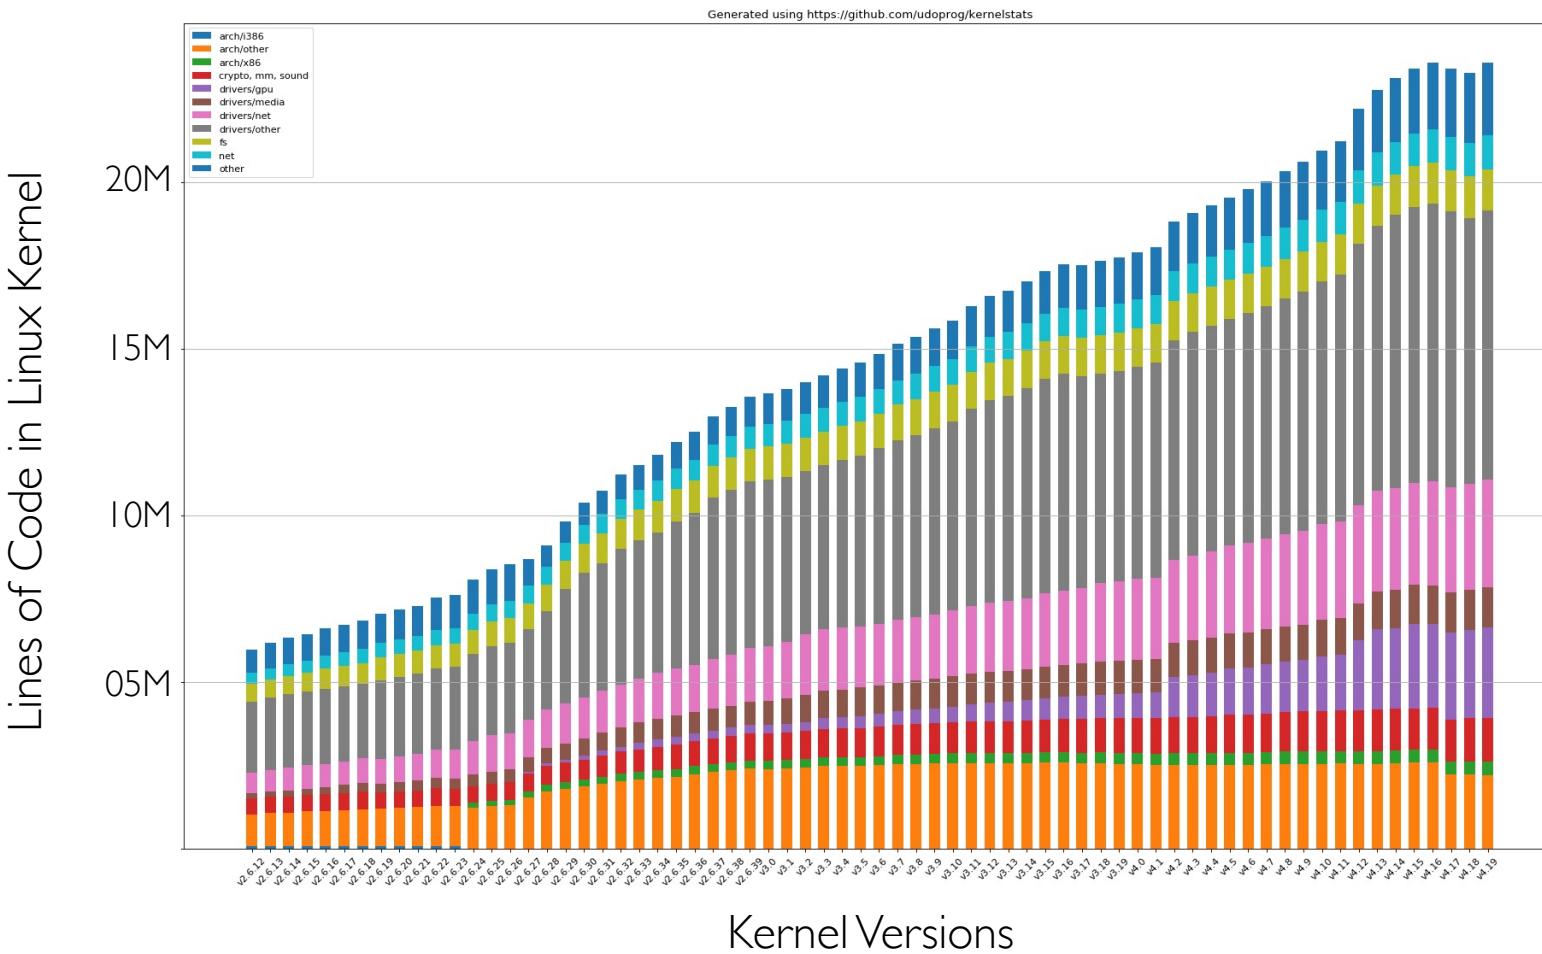
Early vs. Current Operating Systems

- One user/application at any given time
 - Had complete control of hardware
 - OS was runtime library
 - Users would stand in line to use the computer
- **Batch systems**
 - Keep CPU busy by having a queue of jobs
 - OS would load next job while current one runs
 - Users would submit jobs, and wait, and wait, and wait ...
- Multiple users on computer at the same time
 - Multiprogramming: run multiple programs at the same time
 - Interactive performance: try to complete everyone's tasks quickly
 - As computers became cheaper, more important to optimize for user time, not computer time

Complexity

- Applications consisting of...
 - ... a variety of software modules that ...
 - ... run on a variety of devices (machines) that
 - ... implement different hardware architectures
 - ... run competing applications
 - ... fail in unexpected ways
 - ... can be under a variety of attacks
- Not feasible to test software for all possible environments and combinations of components and devices
 - Question is not whether there are bugs but how serious are bugs!

Kernel Complexity



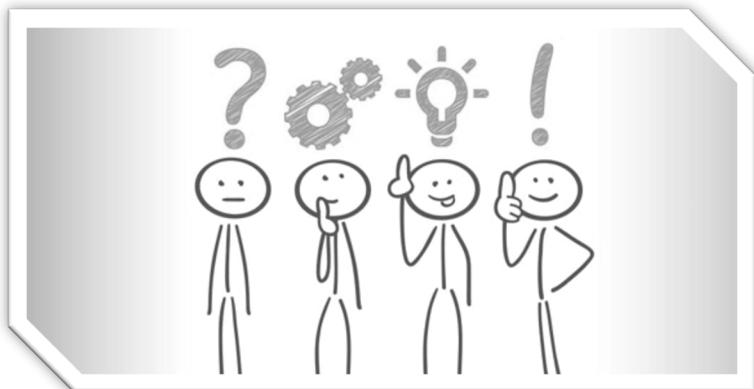
How do We Tame Complexity?

- Every piece of computer hardware different
 - Different CPUs
 - Pentium, ARM, PowerPC, ColdFire
 - Different amounts of memory, disk, ...
 - Different types of devices
 - Mice, keyboards, sensors, cameras, fingerprint readers, touch screen
 - Different networking environment
 - Cable, DSL, Wireless, ...
- Questions
 - Does programmer need to write single program that performs many independent activities?
 - Does every program have to be altered for every piece of hardware?
 - Does one faulty program crash everything?
 - Does every program have access to all hardware?

Summary

- OS provides VM abstraction to handle diverse HW
 - OS simplifies application development by providing standard services
- OS coordinates resources and protect users from each other
 - OS can provide fault containment, fault tolerance, and fault recovery
- ECE 350 combines ideas and concepts from many other areas of computer science and engineering
 - Languages, data structures, hardware, and algorithms

Questions?



Acknowledgment

- Slides by courtesy of Anderson, Culler, Stoica, Silberschatz, Joseph, Canny, and Kumar (Sam)