

# **ECE 350**

## **Real-time**

## **Operating**

## **Systems**



# Lecture 2: OS Concepts

---

Prof. Seyed Majid Zahedi

<https://ece.uwaterloo.ca/~smzahedi>

# Outline

---

- Brief history of OSes
- Four fundamental OS concepts
  - Thread
  - Address space
  - Process
  - Dual-mode operation/protection

# Serial Processing

---

- Machines did not have operating systems
- Run from console with display lights, toggle switches, input device, and printer
- Machine is used by a single user (users had to reserve time to use machines)
- Running programs had long lead time (users had to load compiler and source program, save compiled program, and then load and link it)
- Debugging programs was extremely hard



wikimedia.org



columbia.edu

# Evolution of OSes

---

- Simple batch OS
  - Jobs with same requirement and grouped into batches
  - Special program, called **monitor**, monitors and manages each program
  - Erroneous or misbehaving jobs could corrupt entire system
  - Automatic job sequencing improves throughput, but I/O is still slow
- Multiprogramming batch OS
  - When running job requires I/O, OS switches to another job
  - While this maximizes CPU utilization, response time could still suffer
- Time-sharing OS
  - Multiple users simultaneously access system through terminals
  - Processor's time is shared among multiple users
  - Primary focus is to minimize response time

# Very Brief History of OS

---

- Several distinct phases:
  - Hardware expensive, humans cheap
    - Eniac, ... Multics



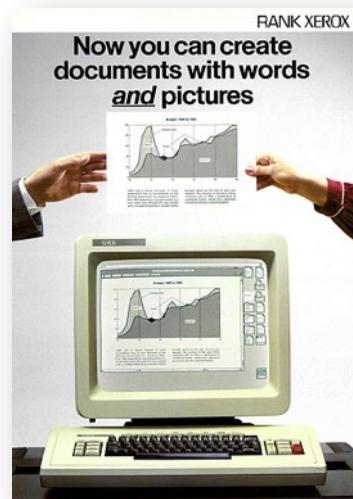
“I think there is a world market for maybe five computers.” – Thomas Watson, chairman of IBM, 1943

Thomas Watson was often called “the world's greatest salesman” by the time of his death in 1956

# Very Brief History of OS (cont.)

---

- Several distinct phases:
  - Hardware expensive, humans cheap
    - Eniac, ... Multics
  - Hardware cheaper, humans expensive
    - PCs, workstations, rise of GUIs
  - Hardware very cheap, humans very expensive
    - Ubiquitous devices, widespread networking



# Very Brief History of OS (cont.)

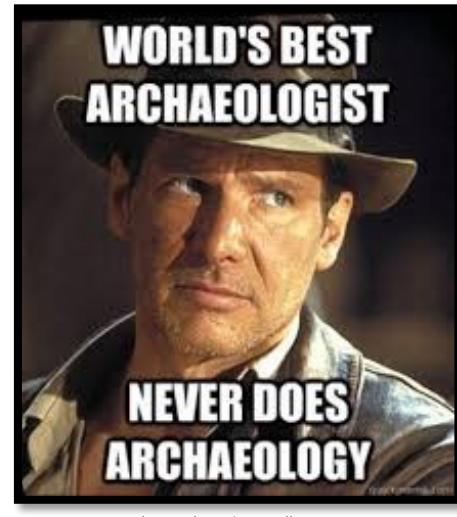
---

- Several distinct phases:
  - Hardware expensive, humans cheap
    - Eniac, ... Multics
  - Hardware cheaper, humans expensive
    - PCs, workstations, rise of GUIs
  - Hardware very cheap, humans very expensive
    - Ubiquitous devices, widespread networking
- Rapid change in hardware leads to changing OS
  - Batch  $\Rightarrow$  multiprogramming  $\Rightarrow$  timesharing  $\Rightarrow$  GUI  $\Rightarrow$  ubiquitous devices
  - Gradual migration of features into smaller machines
- Today
  - Small OS: 100K lines / Large: 20M lines (10M browser!)
  - 100-1000 people-years

# OS Archaeology

---

- Due to high cost of building OS from scratch, most modern OS's have long lineage
- Multics ⇒ AT&T Unix ⇒ BSD Unix ⇒ Ultrix, SunOS, NetBSD,...
- Mach (micro-kernel) + BSD ⇒ NextStep ⇒ XNU ⇒ Apple OS X, iPhone iOS
- MINIX ⇒ Linux ⇒ Android, Chrome OS, RedHat, Ubuntu, Fedora, Debian, Suse,...
- CP/M ⇒ QDOS ⇒ MS-DOS ⇒ Windows 3.1 ⇒ NT ⇒ 95 ⇒ 98 ⇒ 2000 ⇒ XP ⇒ Vista ⇒ 7 ⇒ 8 ⇒ 10 ⇒ ...



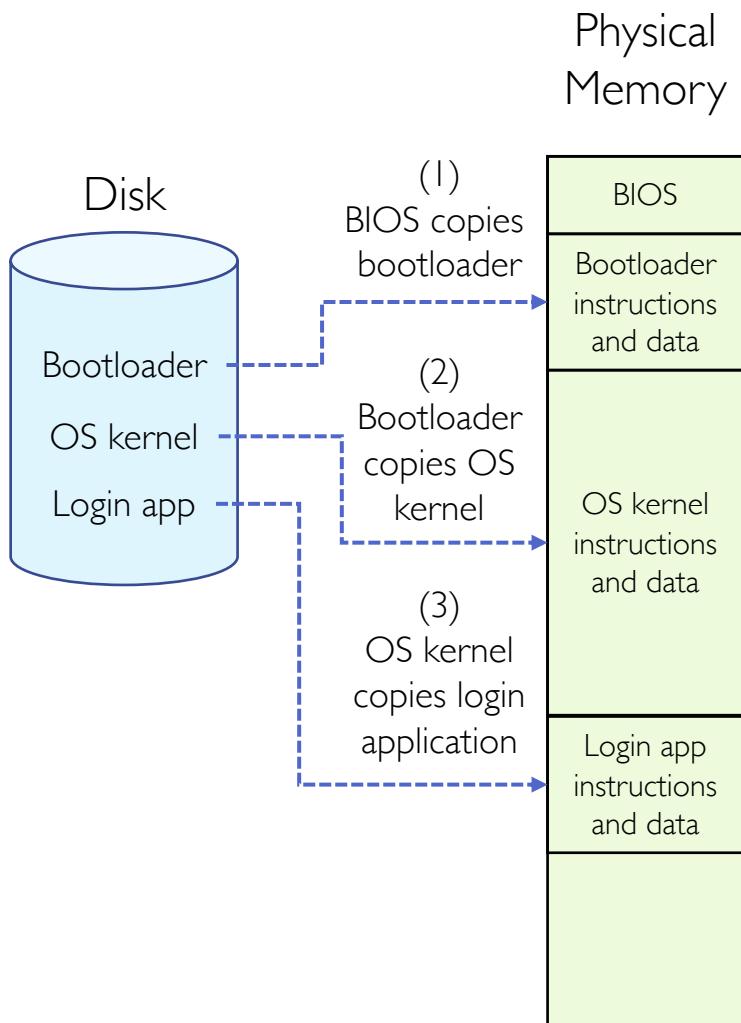
[anthropology4u.medium.com](https://anthropology4u.medium.com)

# Today: Four Fundamental OS Concepts

---

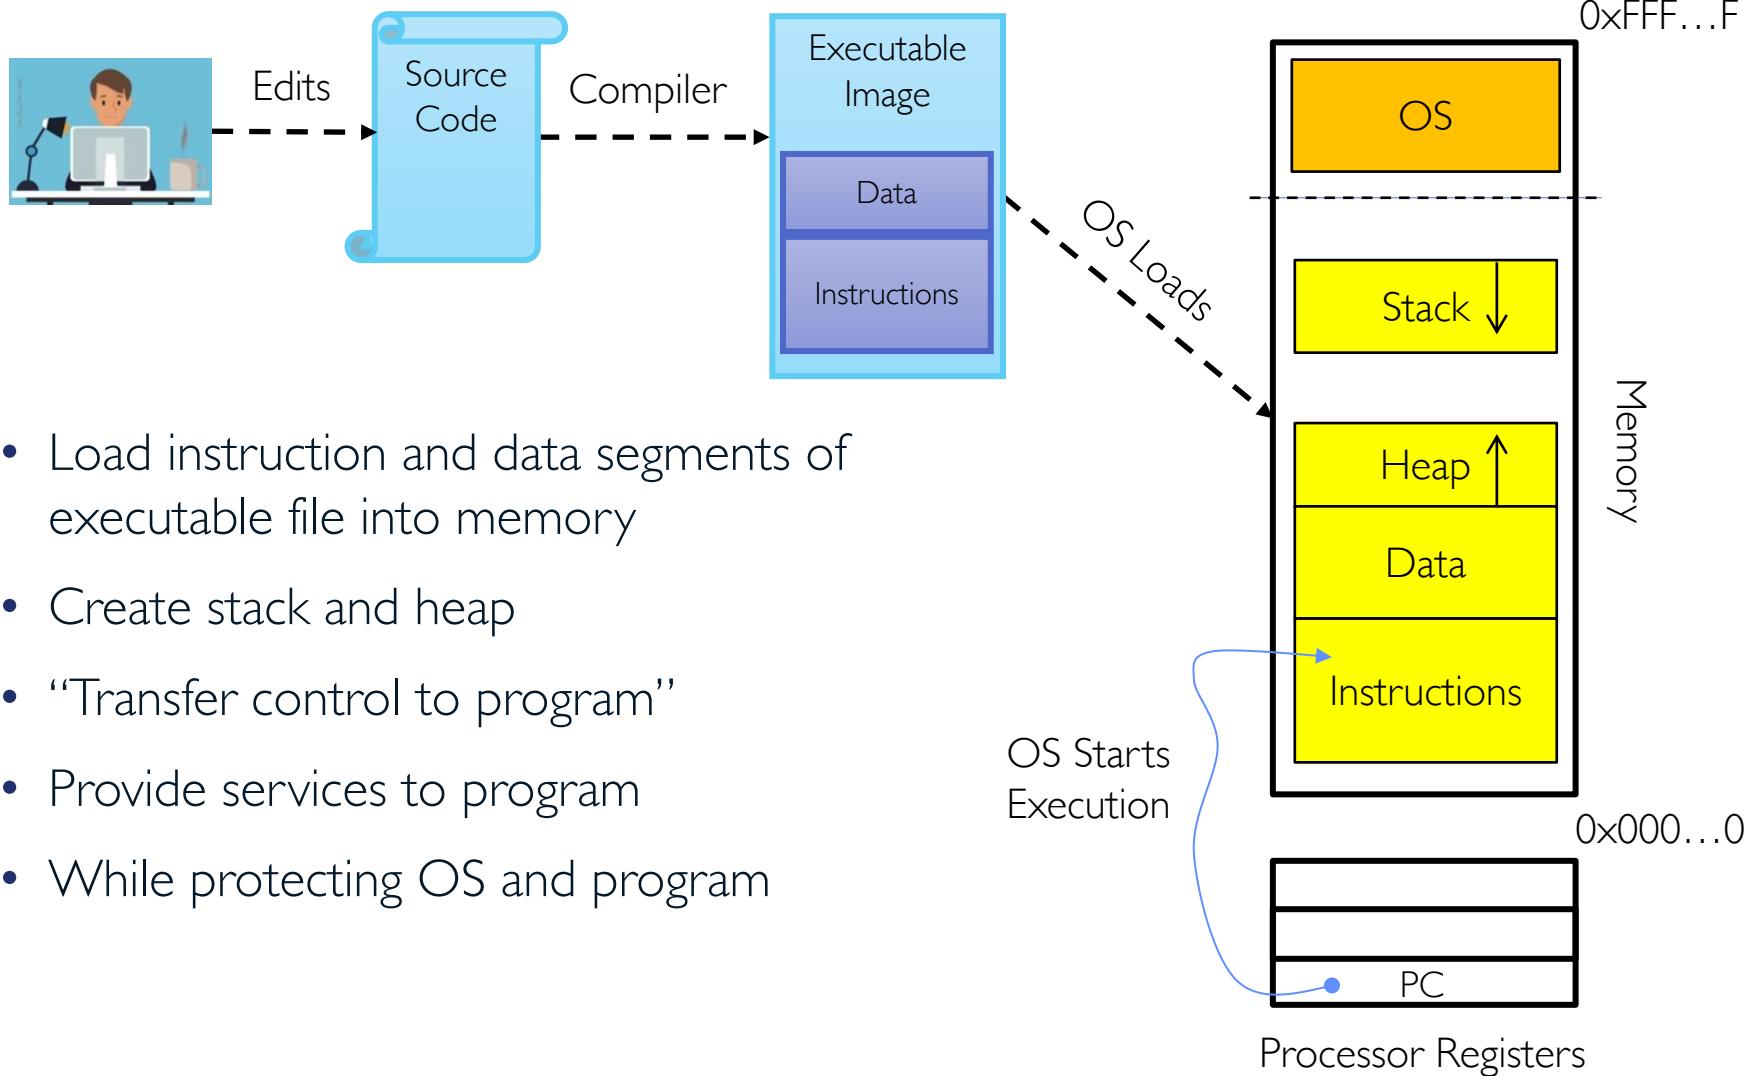
- Thread
  - Single unique execution context which fully describes program state
  - Program counter, registers, execution flags, stack
- Address space (with translation)
  - Address space which is distinct from machine's physical memory addresses
- Process
  - Instance of executing program consisting of address space and 1+ threads
- Dual-mode operation/protection
  - Only "system" can access certain resources
  - OS and hardware are protected from user programs
  - User programs are isolated from one another by controlling translation from program virtual addresses to machine physical addresses

# Booting OS



- In most x86 systems, BIOS is stored on Boot ROM
  - Expensive and writing to it is slow
- Why not storing kernel on Boot ROM?
  - Hard to update (OS updates are frequent)
- Why does BIOS load bootloader not OS?
  - Might have multiple OSes installed
  - BIOS needs to read raw bytes from disk, whereas bootloader needs to know how to read from filesystem

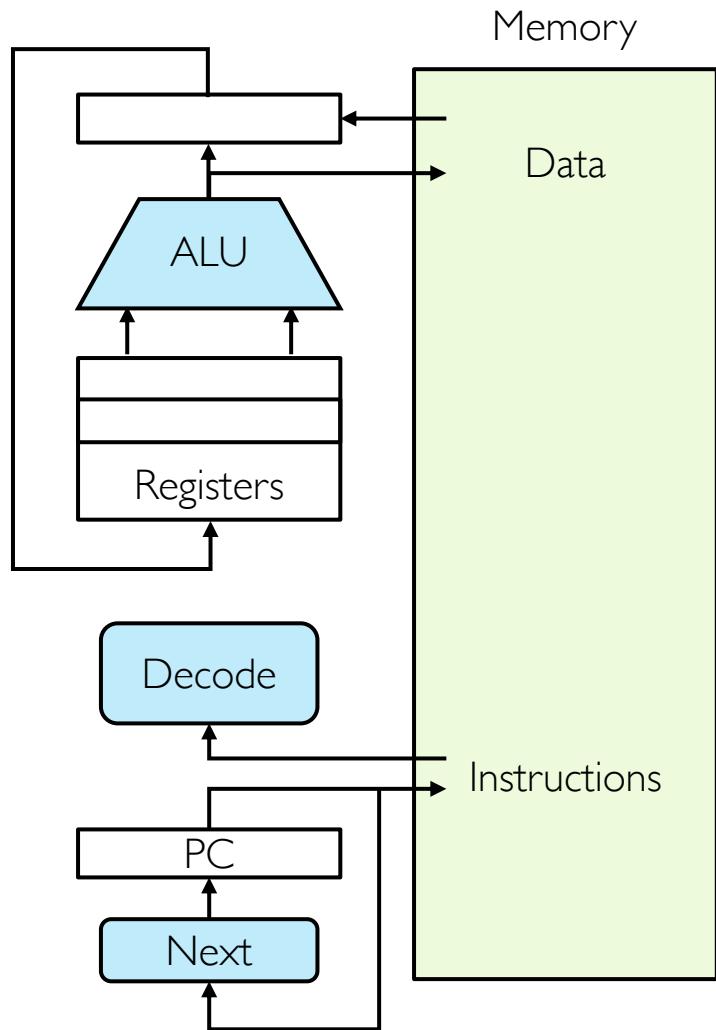
# OS Bottom Line: Run Programs



# Instruction Cycle: Fetch, Decode, Execute

- Execution sequence
  - Fetch instruction at PC
  - Decode
  - Execute (possibly using registers)
  - Write results to registers/memory
  - $PC \leftarrow Next(PC)$
  - Repeat

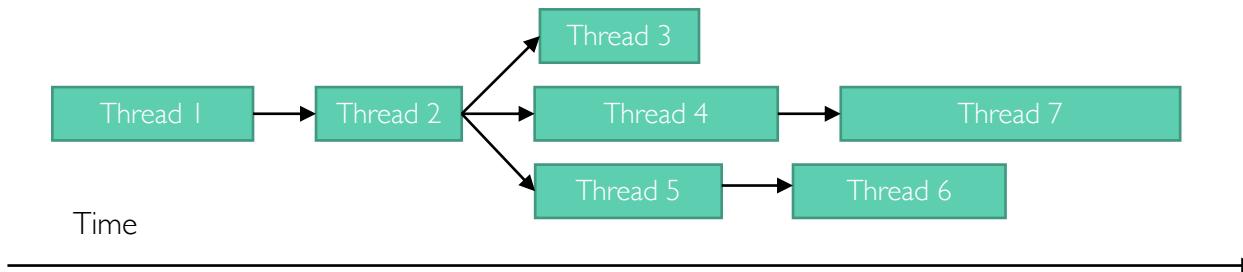
Next instruction or jump  
to new address ...



# Thread (1<sup>st</sup> OS Concept)

---

- Thread is short for **thread of execution**
- Thread of execution is sequence of executable commands that can run on CPU
- Threads have some state and store some local variables
  - Execution state (ready, running, waiting, ...)
  - Saved context when not running
  - Execution stack
  - Local variables
  - ...
- Multithreaded programs use more than one thread (some of the time)
  - Program begins with single initial thread (where the **main** method is)
  - Threads can be created and destroyed within programs dynamically




# Recall: The POSIX Thread

---

- **pthread** refers to POSIX standard that defines thread behavior in UNIX
- **pthread\_create**
  - Creates new thread to run a function
- **pthread\_exit**
  - Quit thread and clean up, wake up joiner if any
  - To allow other threads to continue execution, the main thread should terminate by calling pthread\_exit() rather than exit(3)
- **pthread\_join**
  - In parent, wait for children to exit, then return
- **pthread\_yield**
  - Relinquish CPU voluntarily
- ...

# Recall: Thread Lifecycle

---



# Thread Control Block (TCB)

---


- Data structure in OS containing information needed to manage a thread
  - Thread unique identifier (tid)
  - Stack pointer (points to thread's stack in the process)
  - Program counter (points to the current program instruction of the thread)
  - State of the thread (e.g., running, ready, waiting, etc.)
  - Thread's register values
  - Pointer to process control block (PCB) of the process that the thread lives on

# Address Space (2nd OS Concept)

---

- **Address space**: set of accessible addresses and their state
- **Physical memory**: data storage medium
- **Physical addresses**: addresses available on physical memory
  - For 4GB of memory:  $2^{32}$  ~ 4 billion addresses
- **Virtual addresses**: addresses generated by program
  - For 64-bit processor:  $2^{64} > 18$  quintillion ( $10^{18}$ ) addresses

# Virtual Address Space Layout of C Programs



# Stack Example

---

```
A0: A(int tmp) {  
A1:     if (tmp<2)  
A2:         B();  
A3:         printf(tmp);  
A4:     }  
B0: B() {  
B1:     C();  
B2: }  
C0: C() {  
C1:     A(2);  
C2: }  
ext:  
      A(1);
```

Stack  
Pointer →


```
tmp = 1  
ret = ext
```

- Stack holds temporary results
- Permits recursive execution
- Crucial to modern languages

# Stack Example

---

```
A0: A(int tmp) {  
A1:     if (tmp<2)  
A2:         B();  
A3:         printf(tmp);  
A4:     }  
B0: B() {  
B1:     C();  
B2: }  
C0: C() {  
C1:     A(2);  
C2: }  
        A(1);  
ext:
```



- Stack holds temporary results
- Permits recursive execution
- Crucial to modern languages

# Stack Example

---

```
A0: A(int tmp) {  
A1:     if (tmp<2)  
A2:         B();  
A3:     printf(tmp);  
A4: }  
B0: B() {  
B1:     C();  
B2: }  
C0: C() {  
C1:     A(2);  
C2: }  
        A(1);  
ext:
```

Stack  
Pointer →


```
tmp = 1  
ret = ext
```

- Stack holds temporary results
- Permits recursive execution
- Crucial to modern languages

# Stack Example

---

```
A0: A(int tmp) {  
A1:     if (tmp<2)  
A2:         B();  
A3:         printf(tmp);  
A4:     }  
B0: B() {  
B1:     C();  
B2: }  
C0: C() {  
C1:     A(2);  
C2: }  
        A(1);  
ext:
```




- Stack holds temporary results
- Permits recursive execution
- Crucial to modern languages

# Stack Example

---


```
A0: A(int tmp) {  
A1:     if (tmp<2)  
A2:         B();  
A3:         printf(tmp);  
A4:     }  
  
B0: B() {  
B1:     C();  
B2: }  
  
C0: C() {  
C1:     A(2);  
C2: }  
        A(1);  
  
ext:
```



- Stack holds temporary results
- Permits recursive execution
- Crucial to modern languages

# Stack Example


```
A0: A(int tmp) {  
A1:     if (tmp<2)  
A2:         B();  
A3:         printf(tmp);  
A4:     }  
B0: B() {  
B1:     C();  
B2: }  
C0: C() {  
C1:     A(2);  
C2: }  
        A(1);  
ext:
```



- Stack holds temporary results
- Permits recursive execution
- Crucial to modern languages

# Stack Example

```
A0: A(int tmp) {  
A1:     if (tmp<2)  
A2:         B();  
A3:     printf(tmp);  
A4: }  
B0: B() {  
B1:     C();  
B2: }  
C0: C() {  
C1:     A(2);  
C2: }  
        A(1);  
ext:
```




> 2

- Stack holds temporary results
- Permits recursive execution
- Crucial to modern languages

# Stack Example

```
A0: A(int tmp) {  
A1:     if (tmp<2)  
A2:         B();  
A3:         printf(tmp);  
A4:     }   
B0: B() {  
B1:     C();  
B2: }  
C0: C() {  
C1:     A(2);  
C2: }  
        A(1);  
ext:
```




> 2

- Stack holds temporary results
- Permits recursive execution
- Crucial to modern languages

# Stack Example

---

```
A0: A(int tmp) {  
A1:     if (tmp<2)  
A2:         B();  
A3:         printf(tmp);  
A4:     }  
  
B0: B() {  
B1:     C();  
B2: }  
  
C0: C() {  
C1:     A(2);  
C2: }  
        A(1);  
  
ext:
```




> 2

- Stack holds temporary results
- Permits recursive execution
- Crucial to modern languages

# Stack Example

---

```
A0: A(int tmp) {  
A1:     if (tmp<2)  
A2:         B();  
A3:         printf(tmp);  
A4:     }  
B0: B() {  
B1:     C();  
B2: }  
C0: C() {  
C1:     A(2);  
C2: }  
        A(1);  
ext:
```




> 2

- Stack holds temporary results
- Permits recursive execution
- Crucial to modern languages

# Stack Example

---

```
A0: A(int tmp) {  
A1:     if (tmp<2)  
A2:         B();  
A3:         printf(tmp);  
A4:     }  
B0: B() {  
B1:     C();  
B2: }  
C0: C() {  
C1:     A(2);  
C2: }  
        A(1);  
ext:
```



> 21

- Stack holds temporary results
- Permits recursive execution
- Crucial to modern languages

# Stack Example

---

```
A0: A(int tmp) {  
A1:     if (tmp<2)  
A2:         B();  
A3:         printf(tmp);  
A4:     }   
B0: B() {  
B1:     C();  
B2: }  
C0: C() {  
C1:     A(2);  
C2: }  
        A(1);  
ext:
```

Stack  
Pointer 

```
tmp = 1  
ret = ext
```

> 21

- Stack holds temporary results
- Permits recursive execution
- Crucial to modern languages

# Stack Example


---

```
A0: A(int tmp) {  
A1:     if (tmp<2)  
A2:         B();  
A3:         printf(tmp);  
A4:     }  
B0: B() {  
B1:     C();  
B2: }  
C0: C() {  
C1:     A(2);  
C2: }  
        A(1);  
ext: 
```

> 21

- Stack holds temporary results
- Permits recursive execution
- Crucial to modern languages

# Process (3<sup>rd</sup> OS Concept)




- Process: execution environment with **restricted rights**
  - Address space with one or more threads
  - Owns memory (address space)
  - Owns file descriptors, file system context, ...
- Fundamental tradeoff between **protection and efficiency**
  - Communication easier within a process
  - Communication harder between processes


# Recall: Process Control Block (PCB)

---

- Is a data structure for managing processes
- Is created and updated by OS for each running process
  - Is kept up to date constantly as process executes
- Is held in memory and maintained in some container (e.g., list) by kernel
- Contains everything OS needs to know about the process
  - Unique process identifier (PID), state, priority
  - Program counter (PC)
  - Register data
  - Memory pointers
  - I/O status information,
  - Accounting information
- PC and register data do not need to be updated when program is running
  - They are needed when a system call (trap) or process switch occurs



# Recall: Process Lifecycle (7 States)



As in the five state model, there are some omitted transitions to “finished” and one from “running” to “ready-swapped”

# What Does it Take to Create Process?

---

- Must construct new PCB
  - Inexpensive
- Must set up new *page tables* for address space (more on this later)
  - More expensive
- Copy data from parent process? (Unix `fork()`)
  - With Unix `fork()`, child process gets copy of parent's memory and I/O state
  - Originally very expensive
  - Much less expensive with "copy on write" (more on this later)
- Copy I/O state (file handles, etc.)
  - Medium expense

# Multithreaded Processes

---



- Threads encapsulate **concurrency** and are **active** components
- Address spaces encapsulate **protection** and are **passive** part
  - Keeps buggy program from trashing system
- Why have multiple threads per address space?
  - Processes are expensive to start, switch between, and communicate between

# Multiple Processes vs. Single Process With Multiple Threads


---

- Fundamental tradeoff between protection and efficiency
- Communication harder between processes
  - This is basically IPC
  - It necessarily involves OS
- Communication easier within a process
  - All threads of process share state and resources of process
  - If one thread opens a file, other threads in the process can access it
  - It does not involve OS


# Memory Footprint of Multiple Threads

---


- How do we position stacks relative to each other?
- What maximum size should we choose for stacks?
  - 8KB for **kernel-level** stacks in Linux on x86
  - Less need for tight space constraint for user-level stacks
- What happens if threads violate this?
  - “... program termination and/or corrupted data”
- How might you catch violations?
  - Place guard values at top and bottom of each stack
  - Check values on every context switch



# Multiprogramming: Running Multiple Processes



# Time Sharing: Multiprogramming on Single CPU



- **Illusion:** infinite number of processors
  - Each thread runs on dedicated virtual processor
- **Reality:** few processors, multiple threads running at variable speed
- How can we give illusion of infinite number of processors?
  - **Multiplex in time!**
- How do we **switch** from one process to next?
  - Save PC, SP, and registers in current PCB
  - Load PC, SP, and registers from new PCB
- What **triggers** switch?
  - Timer, voluntary yield, I/O interrupts, ...

# How Do We Multiplex Processes?

---

- **Scheduling:** OS decides which process uses CPU time
  - Only one process is “running” on each CPU at any time
  - Scheduler could give more time to *important* processes
- **Protection:** OS divides non-CPU resources among processes
  - E.g., give each process their own address space
  - E.g., multiplex I/O through system calls

# Scheduling

---

- Kernel scheduler decides which processes/threads receive CPU
- There are variety of scheduling policies for ...
  - Fairness or
  - Realtime guarantees or
  - Latency optimization or ...
- Kernel scheduler maintains data structure containing PCBs




```
if (readyProcesses(PCBs)) {  
    nextPCB = selectProcess(PCBs);  
    run(nextPCB);  
} else {  
    run_idle_process();  
}
```

# Ready Queue

---




- PCBs move from queue to queue as they change state
  - Decisions about which order to remove from queues are **scheduling** decisions
  - Many algorithms possible (more on this in a few weeks)



# Ready Queue And I/O Device Queues

- Process not running  $\Rightarrow$  PCB is in some scheduler queue
  - Separate queue for each device/signal/condition
  - Each queue can have different scheduler policy



# Protection

---




- OS must protect itself from user programs
  - **Reliability**: prevent OS from crashing
  - **Security**: limit scope of what processes can do
  - **Privacy**: limit data each process can access
  - **Fairness**: enforce appropriate share of HW
- It must protect user programs from one another
- Main method is to limit translation from virtual to physical address space

# How to Protect Processes from One Another?


---

- Protection of **memory**
  - Every process does not have access to all memory
- Protection of **I/O devices**
  - Every process does not have access to every device
- Protection of access to **processor**
  - **Preemptive** switching from process to process
  - Use of **timer**
  - Must not be possible to disable timer from user code


# Address Translation Maps: Illusion of Separate Address Space



# Putting it Together: Process




# Putting it Together: Processes




- Switch overhead: **high**
  - CPU state: **low**
  - Memory/IO state: **high**
- Process creation: **high**
- Protection
  - CPU: **yes**
  - Memory/IO: **yes**
- Sharing overhead: **high**  
(involves at least one context switch)

# Putting it Together: Threads




- Switch overhead: **medium**
  - CPU state: **low**
- Thread creation: **medium**
- Protection
  - CPU: **yes**
  - Memory/IO: **no**
- Sharing overhead: **low(ish)**  
(thread switch overhead low)

# Putting it Together: Multi-cores




- Switch overhead: **low** (only CPU state)
- Thread creation: **low**
- Protection
  - CPU: **yes**
  - Memory/IO: **no**
- Sharing overhead: **low** (thread switch overhead **low**, may not need to switch at all!)

# Hyperthreading



- Superscalar processors can execute multiple instructions that are independent
- Multiprocessors can execute multiple independent threads
- Fine-grained multithreading executes two independent threads by switches between them
- Hyperthreading duplicates register state to make second (hardware) “thread” (virtual core)
  - From OS’s point of view, virtual cores are separate CPUs
  - OS can schedule as many threads at a time as there are virtual cores (but, sub-linear speedup!)
  - See: <http://www.cs.washington.edu/research/smt/index.html>

# Putting it Together: Hyperthreading



- Switch overhead between hardware-threads: **very-low** (done in hardware)
- Contention for ALUs/FPUs may **hurt** performance

# Dual-mode Operation (4<sup>th</sup> OS Concept)

---

- Hardware provides at least two modes
  - Kernel mode (or “supervisor” or “protected”)
  - User mode, which is how normal programs are executed
- How can hardware support dual-mode operation?
  - Single bit of state (user/system mode bit)
  - Certain operations/actions only permitted in system/kernel mode
    - In user mode they fail or trap
  - User to kernel transition sets system mode AND saves user PC
    - OS code carefully puts aside user state then performs necessary actions
  - Kernel to user transition clears system mode AND restores user PC
    - E.g., `rfi`: return-from-interrupt

# Three Types of Mode Transfer

---

- **System call**: request for kernel services
  - E.g., **open**, **close**, **read**, **write**, **lseek**
  - Usually implemented by calling *trap* or *syscall* instruction
    - Special instruction is not strictly required; on some systems, processes trigger system calls by executing some instruction with specific invalid opcode
- **Processor exception**: internal, *synchronous*, hardware event
  - E.g., divide by zero, illegal instruction, segmentation fault, page fault
  - Caused by software behavior
- **Interrupt**: external *asynchronous* event
  - E.g., timer, disk ready, network
  - Interrupts can be disabled, exceptions and traps cannot!

# Requirements for Safe Mode Transfer

---


- Limited entry into kernel
  - HW must ensure entry point into kernel is one set up by kernel
  - User programs cannot be allowed to jump to arbitrary locations in kernel
- Atomic changes to processor state
  - In user mode, PC and SP point to memory locations in user process
  - In kernel mode, PC and SP point to memory locations in kernel
  - Mode, PC, SP, and memory protection should all change atomically
- Transparent, restartable execution
  - User-level process could get interrupted between any two instructions
  - OS must restore state of user process exactly as it was before interrupt

# Interrupt Vector Table

- Table set up by OS pointing to code to run on system calls, processor exceptions, and interrupts

- On x86, vector numbers 0-31 are for different types of processor exceptions (e.g., divide-by-zero)

- Vector numbers 32-255 are for different types of interrupts (e.g., timer)
- Vector number 64 is for system call handler



# Interrupt Stack


---

- User process state should be saved
- OS should not save anything on user stack (why?)
  - **Reliability**: what if user program's SP is not valid?
  - **Security**: what if other threads in process change kernel's return address?
- Most OSes go one step further and allocate separate **kernel interrupt stack** (also called **kernel stack**) for each user-level thread
  - PCB could store pointer to kernel stack



memegenerator.net

# Two-stack Model Example



# Interrupt Masking

---


- Interrupt handler runs with interrupts *disabled*
- This simplifies interrupt handling
- Interrupts are re-enabled when interrupt completes
- Interrupts are *deferred* (masked) not ignored
- HW buffers new interrupts until interrupts are re-enabled
- If interrupt are disabled for long time, some interrupts may be lost
- On x86, **cli** disables interrupts and **sti** enables interrupts
  - Only applies to current CPU (on a multicore)
  - User programs cannot use these instructions (why?)

# Mode Transfer Steps in x86

---


- Mask interrupts
- Save PS, SP, and execution flags in temporary HW registers
- Switch onto kernel interrupt stack (specified in special HW register)
- Push the three key values onto interrupt stack
- Optionally save an error code
- Invoke interrupt handler

# Example: x86 Mode Transfer




# Example: x86 Mode Transfer (cont.)

- Single instruction to
  - Save some registers (e.g., SP, PC)
  - Change PC and SP
  - Switch Kernel/user mode



# Example: x86 Mode Transfer (cont.)

- Single instruction to save all registers
- Why is stack pointer saved twice?
  - Hint: are they the same?




# Example: System Call Handler

---


- Vector through well-defined system call entry points!
  - Table mapping system call number to handler
- Locate arguments
  - In registers or on user (!) stack
- Copy arguments (copy before check)
  - From user memory into kernel memory
  - Protect kernel from malicious code evading checks
- Validate arguments
  - Protect kernel from errors in user code
- Copy results back
  - Into user memory

# Basic Cost of System Calls



- Min syscall has  $\sim 25\times$  cost of function call
- Linux vDSO (virtual dynamic shared object) runs some system calls in user space
  - E.g., gettimeofday or getpid

# Aside: Monolithic vs Microkernel OS



# Aside: Influence of Microkernels

---

- Microkernels provide better modularity, security, and fault tolerance, but they introduce higher communication overhead
  - Too many context switches
- Many OSes provide some services externally, like microkernels
  - OS X and Linux: windowing (graphics and UI)
- Some currently monolithic OSes started as microkernels
  - Windows family originally had microkernel design
  - OS X is hybrid of Mach microkernel and FreeBSD monolithic kernel


# Kernel to User Mode Switch Examples

---


- New process/new thread start
  - Jump to first instruction in program/thread
- Return from interrupt, exception, system call
  - Resume suspended execution
- Process/thread context switch
  - Resume some other process
- User-level *upcall* (UNIX *signal*)
  - Asynchronous notification to user program
    - Preemptive user-level threads
    - Asynchronous I/O notification
    - Interprocess communication
    - User-level exception handling
    - User-level resource allocation

# Example: User/Kernel Mode Transfers

---



# System Call Interface: Access Point to Hardware Resources




# Device Drivers

---

- Device-specific code in kernel that interacts directly with device hardware
  - Supports standard, internal interface
  - Same kernel I/O system can interact easily with different device drivers
  - Special device-specific configuration supported with `ioctl()` syscall
- Device drivers are typically divided into two pieces
  - Top half: accessed in call path from system calls
    - implements a set of standard, cross-device calls like `open()`, `close()`, `read()`, `write()`, `ioctl()`, etc.
    - This is kernel's interface to device driver
    - Top half will start I/O to device, may put thread to sleep until finished
  - Bottom half: run as interrupt routine
    - Gets input or transfers next block of output
    - May wake sleeping threads if I/O now complete

# Life Cycle of an I/O Request




# I/O Data Transfer

---


- Programmed I/O
  - Each byte transferred via processor in/out or load/store
  - + Simple hardware, easy to program
  - – Consumes processor cycles proportional to data size
- Direct memory access (DMA)
  - Give controller access to memory bus
  - Ask it to transfer data blocks to/from memory directly

# DMA Transfer



1. Device driver is told to transfer disk data to buffer at address  $x$
2. Device driver tells disk controller to transfer  $C$  bytes from disk to buffer at address  $x$
3. Disk controller initiates DMA transfer
4. Disk controller send each byte to DMA controller
5. DMA controller transfers bytes to buffer  $x$ , increasing address and decreasing  $C$
6. When  $C = 0$ , DMA interrupts CPU to signal transfer completion

# DMA Example: Network Stack in Linux Kernels before 2.6





# How Does Kernel Provide Services?

---




- You said that applications request services from OS via syscall, but ...
  - I've been writing all sorts of applications, and I never ever saw a "syscall" !!!
- That's right!
- It was buried in the programming language runtime library (e.g., libc.a)
  - ... Layering

# OS Run-time Library




# Putting it Together: Web Server

---



# Putting it Together: Web Server (cont.)




# Summary

---

- Thread
  - Single unique execution context which fully describes program state
  - Program counter, registers, execution flags, stack
- Address space (with translation)
  - Address space which is distinct from machine's physical memory addresses
- Process
  - Instance of executing program consisting of address space and 1+ threads
- Dual-mode operation/protection
  - Only "system" can access certain resources
  - OS and hardware are protected from user programs
  - User programs are isolated from one another by controlling translation from program virtual addresses to machine physical addresses

# Questions?

---



# Acknowledgment

---

- Slides by courtesy of Anderson, Culler, Stoica, Silberschatz, Joseph, Zarnett, and Canny