

ECE 350
Real-time
Operating
Systems




Tutorial

Reader/Writer Lock

Prof. Seyed Majid Zahedi

<https://ece.uwaterloo.ca/~smzahedi>

Readers/Writers Lock



- Motivation: consider shared database with two classes of users
 - Readers: never modify database
 - Writers: read and modify database
- Database can have many readers at the same time
- But there can be only one writer active at a time

Properties of Readers/Writers Lock

- Common variant of mutual exclusion
 - One writer at a time, if no readers
 - Many readers, if no writer

Thread 1 \ Thread 2	Writer	Reader
Writer	NO!	NO!
Reader	NO!	OK!

- Correctness constraints
 - Readers can read when no writers
 - Writers can read/write when no readers or writers
 - Only one thread manipulates *state of the lock* at a time

Readers/Writers Lock Class

```
class ReaderWriterLock {
    private:
        Mutex mutex;           // needed to change state vars
        CV okToRead            // CV for readers
        CV okToWrite;         // CV for writers
        int AW = 0;           // # of active writers
        int AR = 0;           // # of active readers
        int WW = 0;           // # of waiting writers
        int WR = 0;           // # of waiting readers

    public:
        void acquireRL();
        void releaseRL();
        void acquireWL();
        void releaseWL();
}
```

Readers/Writers Lock Design Pattern

```
read() {  
    lock.acquireRL();  
  
    // Read shared state  
  
    lock.releaseRL();  
}
```

```
write() {  
    lock.acquireWL();  
  
    // Read/write shared state  
  
    lock.releaseWL();  
}
```

Readers/Writers Lock Implementation

```
acquireRL() {
    mutex.lock(); // Need lock to change state vars
    while (AW + WW > 0) { // Is it safe to read?
        WR++; // No! add to # of waiting readers
        okToRead.wait(&mutex); // Wait on condition variable
        WR--; // No longer waiting
    }
    AR++; // Now we are active again
    mutex.unlock();
}

releaseRL() {
    mutex.lock();
    AR--; // No longer active
    if (AR == 0 && WW > 0) // If no active reader,
        okToWrite.signal(); // wake up waiting writer
    mutex.unlock();
}
```

Why unlock the
mutex here?

Mutex is locked to change internal state of R/W lock

Readers/Writers Lock Implementation (cont.)

```
acquireWL() {
    mutex.lock();
    while (AW + AR > 0) {           // is it safe to write?
        WW++;                       // No! add to # of waiting writers
        okToWrite.wait(&mutex);     // Wait on condition variable
        WW--;                       // No longer waiting
    }
    AW++;                           // Now we are active again
    mutex.unlock();
}
```

```
releaseWL() {
    mutex.lock();
    AW--;                           // No longer active
    if (WW > 0) {                   // Give priority to writers
        okToWrite.signal();         // Wake up a waiting writer
    } else if (WR > 0) {           // If there are waiting readers,
        okToRead.broadcast();       // wake them all up
    }
    mutex.unlock();
}
```


Simulation of Readers/Writers Lock

- Consider following sequence of arrivals: R1, R2, W1, R3

```
read() {
    mutex.lock();
    while (AW + WW > 0) {
        WR++;
        okToRead.wait(&mutex);
        WR--;
    }
    AR++;
    mutex.unlock();

    // Read

    mutex.lock();
    AR--;
    if (AR == 0 && WW > 0)
        okToWrite.signal();
    mutex.unlock();
}
```

```
write() {
    mutex.lock();
    while (AW + AR > 0) {
        WW++;
        okToWrite.wait(&mutex);
        WW--;
    }
    AW++;
    mutex.unlock();

    // Read and Write

    mutex.lock();
    AW--;
    if (WW > 0) {
        okToWrite.signal();
    } else if (WR > 0) {
        okToRead.broadcast();
    }
    mutex.unlock();
}
```

Simulation of Readers/Writers Lock

- RI comes – $AR = 0$, $WR = 0$, $AW = 0$, $WW = 0$

```
read() {  
    mutex.lock();  
    while (AW + WW > 0) {  
        WR++;  
        okToRead.wait(&mutex);  
        WR--;  
    }  
    AR++;  
    mutex.unlock();  
  
    // Read  
  
    mutex.lock();  
    AR--;  
    if (AR == 0 && WW > 0)  
        okToWrite.signal();  
    mutex.unlock();  
}
```

Simulation of Readers/Writers Lock

- RI comes – $AR = 0$, $WR = 0$, $AW = 0$, $WW = 0$

```
read() {  
    mutex.lock();  
    while (AW + WW > 0) {  
        WR++;  
        okToRead.wait(&mutex);  
        WR--;  
    }  
    AR++;  
    mutex.unlock();  
  
    // Read  
  
    mutex.lock();  
    AR--;  
    if (AR == 0 && WW > 0)  
        okToWrite.signal();  
    mutex.unlock();  
}
```

Simulation of Readers/Writers Lock

- RI comes – $AR = 0$, $WR = 0$, $AW = 0$, $WW = 0$

```
read() {
    mutex.lock();
    while (AW + WW > 0) {
        WR++;
        okToRead.wait(&mutex);
        WR--;
    }
    AR++;
    mutex.unlock();

    // Read

    mutex.lock();
    AR--;
    if (AR == 0 && WW > 0)
        okToWrite.signal();
    mutex.unlock();
}
```

Simulation of Readers/Writers Lock

- RI comes – AR = 1, WR = 0, AW = 0, WW = 0

```
read() {
    mutex.lock();
    while (AW + WW > 0) {
        WR++;
        okToRead.wait(&mutex);
        WR--;
    }
    AR++;
    mutex.unlock();

    // Read

    mutex.lock();
    AR--;
    if (AR == 0 && WW > 0)
        okToWrite.signal();
    mutex.unlock();
}
```

Simulation of Readers/Writers Lock

- RI comes – $AR = 1$, $WR = 0$, $AW = 0$, $WW = 0$

```
read() {
    mutex.lock();
    while (AW + WW > 0) {
        WR++;
        okToRead.wait(&mutex);
        WR--;
    }
    AR++;
    mutex.unlock();

    // Read

    mutex.lock();
    AR--;
    if (AR == 0 && WW > 0)
        okToWrite.signal();
    mutex.unlock();
}
```

Simulation of Readers/Writers Lock

- RI comes – $AR = 1$, $WR = 0$, $AW = 0$, $WW = 0$

```
read() {  
    mutex.lock();  
    while (AW + WW > 0) {  
        WR++;  
        okToRead.wait(&mutex);  
        WR--;  
    }  
    AR++;  
    mutex.unlock();  
}
```

```
// Read
```

```
mutex.lock();  
AR--;  
if (AR == 0 && WW > 0)  
    okToWrite.signal();  
mutex.unlock();  
}
```

Simulation of Readers/Writers Lock

- R2 comes – $AR = 1$, $WR = 0$, $AW = 0$, $WW = 0$

```
read() {  
    mutex.lock();  
    while (AW + WW > 0) {  
        WR++;  
        okToRead.wait(&mutex);  
        WR--;  
    }  
    AR++;  
    mutex.unlock();  
  
    // Read  
  
    mutex.lock();  
    AR--;  
    if (AR == 0 && WW > 0)  
        okToWrite.signal();  
    mutex.unlock();  
}
```


Simulation of Readers/Writers Lock

- R2 comes – $AR = 1$, $WR = 0$, $AW = 0$, $WW = 0$

```
read() {  
    mutex.lock();  
    while (AW + WW > 0) {  
        WR++;  
        okToRead.wait(&mutex);  
        WR--;  
    }  
    AR++;  
    mutex.unlock();  
  
    // Read  
  
    mutex.lock();  
    AR--;  
    if (AR == 0 && WW > 0)  
        okToWrite.signal();  
    mutex.unlock();  
}
```

Simulation of Readers/Writers Lock

- R2 comes – $AR = 1$, $WR = 0$, $AW = 0$, $WW = 0$

```
read() {
    mutex.lock();
    while (AW + WW > 0) {
        WR++;
        okToRead.wait(&mutex);
        WR--;
    }
    AR++;
    mutex.unlock();

    // Read

    mutex.lock();
    AR--;
    if (AR == 0 && WW > 0)
        okToWrite.signal();
    mutex.unlock();
}
```

Simulation of Readers/Writers Lock

- R2 comes – AR = 2, WR = 0, AW = 0, WW = 0

```
read() {
    mutex.lock();
    while (AW + WW > 0) {
        WR++;
        okToRead.wait(&mutex);
        WR--;
    }
    AR++;
    mutex.unlock();

    // Read

    mutex.lock();
    AR--;
    if (AR == 0 && WW > 0)
        okToWrite.signal();
    mutex.unlock();
}
```

Simulation of Readers/Writers Lock

- R2 comes – $AR = 2$, $WR = 0$, $AW = 0$, $WW = 0$

```
read() {
    mutex.lock();
    while (AW + WW > 0) {
        WR++;
        okToRead.wait(&mutex);
        WR--;
    }
    AR++;
    mutex.unlock();

    // Read

    mutex.lock();
    AR--;
    if (AR == 0 && WW > 0)
        okToWrite.signal();
    mutex.unlock();
}
```

Simulation of Readers/Writers Lock

- R2 comes – AR = 2, WR = 0, AW = 0, WW = 0

```
read() {  
    mutex.lock();  
    while (AW + WW > 0) {  
        WR++;  
        okToRead.wait(&mutex);  
        WR--;  
    }  
    AR++;  
    mutex.unlock();  
}
```

```
// Read
```

```
mutex.lock();  
AR--;  
if (AR == 0 && WW > 0)  
    okToWrite.signal();  
mutex.unlock();  
}
```

Assume readers take a while to access database
Situation: mutex is unlocked, only AR is non-zero

Simulation of Readers/Writers Lock

- W1 comes – AR = 2, WR = 0, AW = 0, WW = 0

```
write() {  
    mutex.lock();  
    while (AW + AR > 0) {  
        WW++;  
        okToWrite.wait(&mutex);  
        WW--;  
    }  
    AW++;  
    mutex.unlock();  
  
    // Read and Write  
  
    mutex.lock();  
    AW--;  
  
    if (WW > 0) {  
        okToWrite.signal();  
    } else if (WR > 0) {  
        okToRead.broadcast();  
    }  
    mutex.unlock();  
}
```

Simulation of Readers/Writers Lock

- W1 comes – AR = 2, WR = 0, AW = 0, WW = 0

```
write() {  
    mutex.lock();  
    while (AW + AR > 0) {  
        WW++;  
        okToWrite.wait(&mutex);  
        WW--;  
    }  
    AW++;  
    mutex.unlock();  
  
    // Read and Write  
  
    mutex.lock();  
    AW--;  
  
    if (WW > 0) {  
        okToWrite.signal();  
    } else if (WR > 0) {  
        okToRead.broadcast();  
    }  
    mutex.unlock();  
}
```

Simulation of Readers/Writers Lock

- W| comes – AR = 2, WR = 0, AW = 0, WW = 0

```
write() {
    mutex.lock();
    while (AW + AR > 0) {
        WW++;
        okToWrite.wait(&mutex);
        WW--;
    }
    AW++;
    mutex.unlock();

    // Read and Write

    mutex.lock();
    AW--;

    if (WW > 0) {
        okToWrite.signal();
    } else if (WR > 0) {
        okToRead.broadcast();
    }
    mutex.unlock();
}
```


Simulation of Readers/Writers Lock

- W1 comes – AR = 2, WR = 0, AW = 0, WW = 1

```
write() {
    mutex.lock();
    while (AW + AR > 0) {
        WW++;
        okToWrite.wait(&mutex);
        WW--;
    }
    AW++;
    mutex.unlock();

    // Read and Write

    mutex.lock();
    AW--;

    if (WW > 0) {
        okToWrite.signal();
    } else if (WR > 0) {
        okToRead.broadcast();
    }
    mutex.unlock();
}
```

Simulation of Readers/Writers Lock

- W1 comes – AR = 2, WR = 0, AW = 0, WW = 1

W1 cannot start because of readers, so it unlocks mutex and goes to sleep

```
write() {
    mutex.lock();
    while (AW + AR > 0) {
        WW++;
        okToWrite.wait(&mutex);
        WW--;
    }
    AW++;
    mutex.unlock();

    // Read and Write

    mutex.lock();
    AW--;

    if (WW > 0) {
        okToWrite.signal();
    } else if (WR > 0) {
        okToRead.broadcast();
    }
    mutex.unlock();
}
```

Simulation of Readers/Writers Lock

- R3 comes – $AR = 2$, $WR = 0$, $AW = 0$, $WW = 1$

```
read() {  
    mutex.lock();  
    while (AW + WW > 0) {  
        WR++;  
        okToRead.wait(&mutex);  
        WR--;  
    }  
    AR++;  
    mutex.unlock();  
  
    // Read  
  
    mutex.lock();  
    AR--;  
    if (AR == 0 && WW > 0)  
        okToWrite.signal();  
    mutex.unlock();  
}
```

Simulation of Readers/Writers Lock

- R3 comes – $AR = 2$, $WR = 0$, $AW = 0$, $WW = 1$

```
read() {  
    mutex.lock();  
    while (AW + WW > 0) {  
        WR++;  
        okToRead.wait(&mutex);  
        WR--;  
    }  
    AR++;  
    mutex.unlock();  
  
    // Read  
  
    mutex.lock();  
    AR--;  
    if (AR == 0 && WW > 0)  
        okToWrite.signal();  
    mutex.unlock();  
}
```

Simulation of Readers/Writers Lock

- R3 comes – AR = 2, WR = 0, AW = 0, WW = 1

```
read() {
    mutex.lock();
    while (AW + WW > 0) {
        WR++;
        okToRead.wait(&mutex);
        WR--;
    }
    AR++;
    mutex.unlock();

    // Read

    mutex.lock();
    AR--;
    if (AR == 0 && WW > 0)
        okToWrite.signal();
    mutex.unlock();
}
```

Simulation of Readers/Writers Lock

- R3 comes – AR = 2, WR = 1, AW = 0, WW = 1

```
read() {
    mutex.lock();
    while (AW + WW > 0) {
        WR++;
        okToRead.wait(&mutex);
        WR--;
    }
    AR++;
    mutex.unlock();

    // Read

    mutex.lock();
    AR--;
    if (AR == 0 && WW > 0)
        okToWrite.signal();
    mutex.unlock();
}
```

Simulation of Readers/Writers Lock

- R3 comes – AR = 2, WR = 1, AW = 0, WW = 1

```
read() {
    mutex.lock();
    while (AW + WW > 0) {
        WR++;
        okToRead.wait(&mutex);
        WR--;
    }
    AR++;
    mutex.unlock();

    // Read

    mutex.lock();
    AR--;
    if (AR == 0 && WW > 0)
        okToWrite.signal();
    mutex.unlock();
}
```

Status:

- R1 and R2 still reading
- W1 and R3 waiting on `okToWrite` and `okToRead`, respectively

Simulation of Readers/Writers Lock

- R2 is done reading – $AR = 2$, $WR = 1$, $AW = 0$, $WW = 1$

```
read() {
    mutex.lock();
    while (AW + WW > 0) {
        WR++;
        okToRead.wait(&mutex);
        WR--;
    }
    AR++;
    mutex.unlock();

    // Read

    mutex.lock();
    AR--;
    if (AR == 0 && WW > 0)
        okToWrite.signal();
    mutex.unlock();
}
```


Simulation of Readers/Writers Lock

- R2 is done reading – $AR = 1$, $WR = 1$, $AW = 0$, $WW = 1$

```
read() {
    mutex.lock();
    while (AW + WW > 0) {
        WR++;
        okToRead.wait(&mutex);
        WR--;
    }
    AR++;
    mutex.unlock();

    // Read

    mutex.lock():
    AR--;
    if (AR == 0 && WW > 0)
        okToWrite.signal();
    mutex.unlock();
}
```

Simulation of Readers/Writers Lock

- R2 is done reading – $AR = 1$, $WR = 1$, $AW = 0$, $WW = 1$

```
read() {
    mutex.lock();
    while (AW + WW > 0) {
        WR++;
        okToRead.wait(&mutex);
        WR--;
    }
    AR++;
    mutex.unlock();

    // Read

    mutex.lock();
    AR--;
    if (AR == 0 && WW > 0)
        okToWrite.signal();
    mutex.unlock();
}
```

Simulation of Readers/Writers Lock

- R2 is done reading – $AR = 1$, $WR = 1$, $AW = 0$, $WW = 1$

```
read() {
    mutex.lock();
    while (AW + WW > 0) {
        WR++;
        okToRead.wait(&mutex);
        WR--;
    }
    AR++;
    mutex.unlock();

    // Read

    mutex.lock();
    AR--;
    if (AR == 0 && WW > 0)
        okToWrite.signal();
    mutex.unlock();
}
```

Simulation of Readers/Writers Lock

- RI is done reading – $AR = 1$, $WR = 1$, $AW = 0$, $WW = 1$

```
read() {
    mutex.lock();
    while (AW + WW > 0) {
        WR++;
        okToRead.wait(&mutex);
        WR--;
    }
    AR++;
    mutex.unlock();

    // Read

    mutex.lock();
    AR--;
    if (AR == 0 && WW > 0)
        okToWrite.signal();
    mutex.unlock();
}
```

Simulation of Readers/Writers Lock

- RI is done reading – $AR = 0$, $WR = 1$, $AW = 0$, $WW = 1$

```
read() {
    mutex.lock();
    while (AW + WW > 0) {
        WR++;
        okToRead.wait(&mutex);
        WR--;
    }
    AR++;
    mutex.unlock();

    // Read

    mutex.lock():
    AR--;
    if (AR == 0 && WW > 0)
        okToWrite.signal();
    mutex.unlock();
}
```

Simulation of Readers/Writers Lock

- RI is done reading – $AR = 0$, $WR = 1$, $AW = 0$, $WW = 1$

```
read() {
    mutex.lock();
    while (AW + WW > 0) {
        WR++;
        okToRead.wait(&mutex);
        WR--;
    }
    AR++;
    mutex.unlock();

    // Read

    mutex.lock();
    AR--;
    if (AR == 0 && WW > 0)
        okToWrite.signal();
    mutex.unlock();
}
```

Simulation of Readers/Writers Lock

- R1 is done reading – $AR = 0$, $WR = 1$, $AW = 0$, $WW = 1$

```
read() {
    mutex.lock();
    while (AW + WW > 0) {
        WR++;
        okToRead.wait(&mutex);
        WR--;
    }
    AR++;
    mutex.unlock();

    // Read

    mutex.lock();
    AR--;
    if (AR == 0 && WW > 0)
        okToWrite.signal();
    mutex.unlock();
}
```

All active readers are finished, R1 signals waiting writer – note, R3 is still waiting

Simulation of Readers/Writers Lock

- RI is done reading – $AR = 0$, $WR = 1$, $AW = 0$, $WW = 1$

```
read() {
    mutex.lock();
    while (AW + WW > 0) {
        WR++;
        okToRead.wait(&mutex);
        WR--;
    }
    AR++;
    mutex.unlock();

    // Read

    mutex.lock();
    AR--;
    if (AR == 0 && WW > 0)
        okToWrite.signal();
    mutex.unlock();
}
```


Simulation of Readers/Writers Lock

- WI gets a signal – $AR = 0$, $WR = 1$, $AW = 0$, $WW = 1$

WI gets signal from RI

```
write() {
    mutex.lock();
    while (AW + AR > 0) {
        WW++;
        okToWrite.wait(&mutex);
        WW--;
    }
    AW++;
    mutex.unlock();

    // Read and Write

    mutex.lock();
    AW--;

    if (WW > 0) {
        okToWrite.signal();
    } else if (WR > 0) {
        okToRead.broadcast();
    }
    mutex.unlock();
}
```

Simulation of Readers/Writers Lock

- W| gets a signal – $AR = 0$, $WR = 1$, $AW = 0$, $WW = 0$

```
write() {
    mutex.lock();
    while (AW + AR > 0) {
        WW++;
        okToWrite.wait(&mutex);
    }
    AW++;
    mutex.unlock();

    // Read and Write

    mutex.lock();
    AW--;

    if (WW > 0) {
        okToWrite.signal();
    } else if (WR > 0) {
        okToRead.broadcast();
    }
    mutex.unlock();
}
```

Simulation of Readers/Writers Lock

- W| gets a signal – $AR = 0$, $WR = 1$, $AW = 1$, $WW = 0$

```
write() {
    mutex.lock();
    while (AW + AR > 0) {
        WW++;
        okToWrite.wait(&mutex);
        WW--;
    }
    AW++;
    mutex.unlock();

    // Read and Write

    mutex.lock();
    AW--;

    if (WW > 0) {
        okToWrite.signal();
    } else if (WR > 0) {
        okToRead.broadcast();
    }
    mutex.unlock();
}
```

Simulation of Readers/Writers Lock

- W| gets a signal – $AR = 0$, $WR = 1$, $AW = 1$, $WW = 0$

```
write() {
    mutex.lock();
    while (AW + AR > 0) {
        WW++;
        okToWrite.wait(&mutex);
        WW--;
    }
    AW++;
    mutex.unlock();

    // Read and Write

    mutex.lock();
    AW--;

    if (WW > 0) {
        okToWrite.signal();
    } else if (WR > 0) {
        okToRead.broadcast();
    }
    mutex.unlock();
}
```

Simulation of Readers/Writers Lock

- W| gets a signal – $AR = 0$, $WR = 1$, $AW = 1$, $WW = 0$

```
write() {
    mutex.lock();
    while (AW + AR > 0) {
        WW++;
        okToWrite.wait(&mutex);
        WW--;
    }
    AW++;
    mutex.unlock();

    // Read and Write

    mutex.lock();
    AW--;

    if (WW > 0) {
        okToWrite.signal();
    } else if (WR > 0) {
        okToRead.broadcast();
    }
    mutex.unlock();
}
```

Simulation of Readers/Writers Lock

- W| is done – $AR = 0$, $WR = 1$, $AW = 1$, $WW = 0$

```
write() {
    mutex.lock();
    while (AW + AR > 0) {
        WW++;
        okToWrite.wait(&mutex);
        WW--;
    }
    AW++;
    mutex.unlock();

    // Read and Write

    mutex.lock();
    AW--;

    if (WW > 0) {
        okToWrite.signal();
    } else if (WR > 0) {
        okToRead.broadcast();
    }
    mutex.unlock();
}
```

Simulation of Readers/Writers Lock

- WI is done – AR = 0, WR = 1, AW = 0, WW = 0

```
write() {
    mutex.lock();
    while (AW + AR > 0) {
        WW++;
        okToWrite.wait(&mutex);
        WW--;
    }
    AW++;
    mutex.unlock();

    // Read and Write

    mutex.lock();
    AW--;

    if (WW > 0) {
        okToWrite.signal();
    } else if (WR > 0) {
        okToRead.broadcast();
    }
    mutex.unlock();
}
```

Simulation of Readers/Writers Lock

- WI is done – $AR = 0$, $WR = 1$, $AW = 0$, $WW = 0$

```
write() {
    mutex.lock();
    while (AW + AR > 0) {
        WW++;
        okToWrite.wait(&mutex);
        WW--;
    }
    AW++;
    mutex.unlock();

    // Read and Write

    mutex.lock();
    AW--;

    if (WW > 0) {
        okToWrite.signal();
    } else if (WR > 0) {
        okToRead.broadcast();
    }
    mutex.unlock();
}
```


Simulation of Readers/Writers Lock

- W1 is done – AR = 0, WR = 1, AW = 0, WW = 0

```
write() {
    mutex.lock();
    while (AW + AR > 0) {
        WW++;
        okToWrite.wait(&mutex);
        WW--;
    }
    AW++;
    mutex.unlock();

    // Read and Write

    mutex.lock();
    AW--;

    if (WW > 0) {
        okToWrite.signal();
    } else if (WR > 0) {
        okToRead.broadcast();
    }
    mutex.unlock();
}
```

Simulation of Readers/Writers Lock

- W1 is done – $AR = 0$, $WR = 1$, $AW = 0$, $WW = 0$

```
write() {
    mutex.lock();
    while (AW + AR > 0) {
        WW++;
        okToWrite.wait(&mutex);
        WW--;
    }
    AW++;
    mutex.unlock();

    // Read and Write

    mutex.lock();
    AW--;

    if (WW > 0) {
        okToWrite.signal();
    } else if (WR > 0) {
        okToRead.broadcast();
    }
    mutex.unlock();
}
```

No waiting writer, so
only signal R3

Simulation of Readers/Writers Lock

- W1 is done – AR = 0, WR = 1, AW = 0, WW = 0

```
write() {
    mutex.lock();
    while (AW + AR > 0) {
        WW++;
        okToWrite.wait(&mutex);
        WW--;
    }
    AW++;
    mutex.unlock();

    // Read and Write

    mutex.lock();
    AW--;

    if (WW > 0) {
        okToWrite.signal();
    } else if (WR > 0) {
        okToRead.broadcast();
    }
    mutex.unlock();
}
```

Simulation of Readers/Writers Lock

- R3 gets a signal – $AR = 0$, $WR = 1$, $AW = 0$, $WW = 0$

```
read() {  
    mutex.lock();  
    while (AW + WW > 0) {  
        WR++;  
        okToRead.wait(&mutex);  
        WR--;  
    }  
    AR++;  
    mutex.unlock();  
  
    // Read  
  
    mutex.lock();  
    AR--;  
    if (AR == 0 && WW > 0)  
        okToWrite.signal();  
    mutex.unlock();  
}
```

R3 gets signal from W3

Simulation of Readers/Writers Lock

- R3 gets a signal – $AR = 0$, $WR = 0$, $AW = 0$, $WW = 0$

```
read() {
    mutex.lock();
    while (AW + WW > 0) {
        WR++;
        okToRead.wait(&mutex);
        WR--;
    }
    AR++;
    mutex.unlock();

    // Read

    mutex.lock();
    AR--;
    if (AR == 0 && WW > 0)
        okToWrite.signal();
    mutex.unlock();
}
```

Simulation of Readers/Writers Lock

- R3 gets a signal – AR = 1, WR = 0, AW = 0, WW = 0

```
read() {  
    mutex.lock();  
    while (AW + WW > 0) {  
        WR++;  
        okToRead.wait(&mutex);  
        WR--;  
    }  
    AR++;  
    mutex.unlock();  
  
    // Read  
  
    mutex.lock();  
    AR--;  
    if (AR == 0 && WW > 0)  
        okToWrite.signal();  
    mutex.unlock();  
}
```

Simulation of Readers/Writers Lock

- R3 finishes – $AR = 0$, $WR = 0$, $AW = 0$, $WW = 0$

```
read() {
    mutex.lock();
    while (AW + WW > 0) {
        WR++;
        okToRead.wait(&mutex);
        WR--;
    }
    AR++;
    mutex.unlock();

    // Read

    mutex.lock();
    AR--;
    if (AR == 0 && WW > 0)
        okToWrite.signal();
    mutex.unlock();
}
```

Readers/Writers Lock Questions

```
read() {
    mutex.lock();
    while (AW + WW > 0) {
        WR++;
        okToRead.wait(&mutex);
        WR--;
    }
    AR++;
    mutex.unlock();

    // Read

    mutex.lock();
    AR--;
    if (AR == 0 && WW > 0)
        okToWrite.signal();
    mutex.unlock();
}
```

What if we
remove this line?

It works but it's inefficient, writer wakes up and goes to sleep again when it's not safe to write

```
write() {
    mutex.lock();
    while (AW + AR > 0) {
        WW++;
        okToWrite.wait(&mutex);
        WW--;
    }
    AW++;
    mutex.unlock();

    // Read and Write

    mutex.lock();
    AW--;

    if (WW > 0) {
        okToWrite.signal();
    } else if (WR > 0) {
        okToRead.broadcast();
    }
    mutex.unlock();
}
```


Readers/Writers Lock Questions

```
read() {
    mutex.lock();
    while (AW + WW > 0) {
        WR++;
        okToRead.wait(&mutex);
        WR--;
    }
    AR++;
    mutex.unlock();

    // Read

    mutex.lock();
    AR--;
    if (AR == 0 && WW > 0)
        okToWrite.broadcast();
    mutex.unlock();
}
```

What if we turn
signal() to
broadcast()?

```
write() {
    mutex.lock();
    while (AW + AR > 0) {
        WW++;
        okToWrite.wait(&mutex);
        WW--;
    }
    AW++;
    mutex.unlock();

    // Read and Write

    mutex.lock();
    AW--;

    if (WW > 0) {
        okToWrite.signal();
    } else if (WR > 0) {
        okToRead.broadcast();
    }
    mutex.unlock();
}
```

It works but it's inefficient to wake up all writers
only for one to becomes active

Readers/Writers Lock Questions

```
read() {
    mutex.lock();
    while (AW + WW > 0) {
        WR++;
        okToContinue.wait(&mutex);
        WR--;
    }
    AR++;
    mutex.unlock();
```

// Read

```
    mutex.lock();
    AR--;
    if (AR == 0 && WW > 0)
        okToContinue.signal();
    mutex.unlock();
}
```

What if we turn `okToWrite` and `okToRead` into `okContinue`?

Signal could be delivered to wrong thread (reader) and get waived!

```
write() {
    mutex.lock();
    while (AW + AR > 0) {
        WW++;
        okToContinue.wait(&mutex);
        WW--;
    }
    AW++;
    mutex.unlock();
```

// Read and Write

```
    mutex.lock();
    AW--;
    if (WW > 0) {
        okToContinue.signal();
    } else if (WR > 0) {
        okToContinue.broadcast();
    }
    mutex.unlock();
}
```

Readers/Writers Lock Questions

```
read() {
    mutex.lock();
    while (AW + WW > 0) {
        WR++;
        okToContinue.wait(&mutex);
        WR--;
    }
    AR++;
    mutex.unlock();

    // Read

    mutex.lock();
    AR--;
    if (AR == 0 && WW > 0)
        okToContinue.broadcast();
    mutex.unlock();
}
```

Does changing `signal()` to `broadcast()` solve the problem?

Yes, but it's inefficient to wake up all threads for only one to become active

```
write() {
    mutex.lock();
    while (AW + AR > 0) {
        WW++;
        okToContinue.wait(&mutex);
        WW--;
    }
    AW++;
    mutex.unlock();

    // Read and Write

    mutex.lock();
    AW--;

    if (WW > 0) {
        okToContinue.broadcast();
    } else if (WR > 0) {
        okToContinue.broadcast();
    }
    mutex.unlock();
}
```

Readers/Writers Lock Questions

```
read() {
    mutex.lock();
    while (AW + WW > 0) {
        WR++;
        okToRead.wait(&mutex);
        WR--;
    }
    AR++;
    mutex.unlock();
```

// Read

```
    mutex.lock();
    AR--;
    if (AR == 0 && WW > 0)
        okToWrite.signal();
    mutex.unlock();
}
```

Can readers starve?

Yes: writers take priority

```
write() {
    mutex.lock();
    while (AW + AR > 0) {
        WW++;
        okToWrite.wait(&mutex);
        WW--;
    }
    AW++;
    mutex.unlock();
```

// Read and Write

```
    mutex.lock();
    AW--;

    if (WW > 0) {
        okToWrite.signal();
    } else if (WR > 0) {
        okToRead.broadcast();
    }
    mutex.unlock();
}
```

Readers/Writers Lock Questions

```
read() {
    mutex.lock();
    while (AW + WW > 0) {
        WR++;
        okToRead.wait(&mutex);
        WR--;
    }
    AR++;
    mutex.unlock();
```

// Read

```
mutex.lock();
AR--;
if (AR == 0 && WW > 0)
    okToWrite.signal();
mutex.unlock();
}
```

Can writers starve?

Yes: a waiting writer may not be able to proceed if another writer slips in between signal and wakeup

```
write() {
    mutex.lock();
    while (AW + AR > 0) {
        WW++;
        okToWrite.wait(&mutex);
        WW--;
    }
    AW++;
    mutex.unlock();
```

// Read and Write

```
mutex.lock();
AW--;
if (WW > 0) {
    okToWrite.signal();
} else if (WR > 0) {
    okToRead.broadcast();
}
mutex.unlock();
}
```

Readers/Writers Lock Without Writer Starvation (Take I)

check also for waiting writers

```
acquireWL() {  
    mutex.lock();  
    while (AW + AR + WW > 0) {  
        WW++;  
        okToWrite.wait(&mutex);  
        WW--;  
    }  
    AW++;  
    mutex.unlock();  
}
```

- Does this work?
 - No! If there **WW** is more than zero, then no waiting writer can successfully proceed

Readers/Writers Lock Without Writer Starvation (Take 2)

Idea: keep track of writers' waiting order, allow writer with longest waiting time to proceed

Does this work?

```
numWriters = 0;  
nextToGo = 1;
```

```
acquireWL() {  
    mutex.lock();  
    myPos = numWriters++;  
    while (AW + AR > 0 ||  
           myPos > nextToGo) {  
        WW++;  
        okToWrite.wait(&mutex);  
        WW--;  
    }  
    AW++;  
    mutex.unlock();  
}
```

No, Signal can wake up a wrong writer and get waived!

```
releaseWL() {  
    mutex.lock();  
    AW--;  
    nextToGo++;  
    if (WW > 0) {  
        okToWrite.signal();  
    } else if (WR > 0) {  
        okToRead.broadcast();  
    }  
    mutex.unlock();  
}
```

Inefficient solution is to change `signal()` to `broadcast()`

Readers/Writers Lock Without Writer Starvation (Take 3)

Idea for efficient solution: have separate CV for each waiting writer and put CV's in ordered queue


```
numWriters = 0;
nextToGo = 1;

acquireWL() {
    mutex.lock();
    myPos = numWriters++;
    myCV = new CV();
    queue.enqueue(myCV);
    while (AW + AR > 0 ||
           myPos > nextToGo) {
        WW++;
        myCV.wait(&mutex);
        WW--;
    }
    AW++;
    queue.dequeue();
    mutex.unlock();
}
```

```
releaseWL() {
    mutex.lock();
    AW--;
    nextToGo++;
    if (WW > 0) {
        queue.head().signal();
    } else if (WR > 0) {
        okToRead.broadcast();
    }
    mutex.unlock();
}

releaseRL() {
    mutex.lock();
    AR--;
    if (AR == 0 && WW > 0)
        queue.head().signal();
    mutex.unlock();
}
```


Questions?



Acknowledgment

- Slides by courtesy of Anderson, Culler, Stoica, Silberschatz, Joseph, and Canny